



Ciencia de Datos

2024

Introducción a la Ciencia de Datos


OPTATIVA - LICENCIATURA EN INFORMÁTICA
FACET-UNT

Ciencia de Datos: Fundamentos y Herramientas



CURSO DE POSTGRADO - FACET-UNT



Deep Learning



Extraer patrones de los
datos usando redes
neuronales







- <https://this-person-does-not-exist.com/en>

Recurrent Neural Networks

Redes neuronales recurrentes

- Para procesar secuencias de datos $x(t) = x(1), \dots, x(\tau)$
- Recurrente -> ejecuta la misma tarea para cada elemento de la secuencia, la salida depende de varias cuentas realizadas en muestras previas
- Tiene “memoria”




Recurrent Neural Networks

RNNs como una manera de modelar problemas de secuencias de datos

Lo que necesitaría el modelo es:

- Manejar secuencias de **longitud variable**
- Trackear dependencias de **largo plazo (long-term dependences)**
- Mantener la información sobre el **orden**
- **Compartir parámetros** a lo largo de la secuencia

Queremos pronosticar esto!







Recurrent Neural Networks

Los RNN es un approach to sequence modeling problems
que no necesitan el orden.

- Manejar secuencias de **longitud variable**
- Trackear dependencias de **largo plazo (long-term dependences)**
- Mantener la información sobre el **orden**
- **Compartir parámetros** a lo largo de la secuencia

Cuantos pasos? (longitud)



Y + arquitecturas


Recurrent Neural Networks

RNNs as an approach to sequence modeling problems

Lo que necesitamos el modelo:

- Manejar secuencias de **longitud variable**
- Trackear dependencias de **largo plazo (long-term dependences)**
- Mantener la informacion sobre el **orden**
- **Compartir parametros** a lo largo de la secuencia

Que tan importante es la informacion del pasado (lejano) ?




La serie podria tener comportamientos regulares (diarios, estacionales, etc.)

Recurrent Neural Networks

RNNs as an approach to sequence modeling problems

- Manejar secuencias de **longitud variable**
- Trackear dependencias de **largo plazo (long-term dependences)**
- Mantener la informacion sobre el **orden**
- **Compartir parametros** a lo largo de la secuencia




$vTEC(t-2), vTEC(t-1), vTEC(t-0) \leftrightarrow vTEC(t-0), vTEC(t-2), vTEC(t-1)$

Recurrent Neural Networks



RNNs as an approach to sequence modeling problems
Lo que necesitan el modelo:

- Manejar secuencias de **longitud variable**
- Trackear dependencias de **largo plazo (long-term dependences)**
- Mantener la informacion sobre el **orden**
- **Compartir parametros** a lo largo de la secuencia



RNNs tienen un estado o **state** (h_t), que es **actualizado** en cada **paso** a medida que la secuencia es procesada usando los **mismos parametros** en cada paso de tiempo

Recurrent Neural Networks




$$h_t = f_W(h_{t-1}, x_t)$$



cell state function old state input vector at
parameterized by W time step t

- Aplica una **relación recurrente** en cada paso de tiempo para procesar la secuencia

Recurrent Neural Networks




Backpropagation through time: long time dependences



- Toma el gradiente de la perdida (loss) con respecto a cada parametro
- Cambia parametros para poder minimizar la perdida


Backpropagation through time: long time dependences



- Realiza la computacion individual de l_i para cada paso y los suma
- Hace backpropagation de los errores individuales de cada paso y luego de todos los pasos de tiempo hasta el inicio de la secuencia

<https://kharshit.github.io/blog/2019/02/22/backpropagation – through – time>

Backpropagation through time: long time dependences



- Mucho tiempo de computo!

- Muchos valores $\gg 1$ -> exploding gradient (*)
- Muchos valores $\ll 1$ -> vanishing gradient

(*) Gradient clipping is a simple technique: If the gradient gets too large, we rescale it to keep it small.

Vanishing gradient problem



Multiplica **muchos numeros pequeños**



Errores debido a que para aquellos pasos lejanos en el tiempo, los gradientes se hacen cada vez mas pequeños



Sesga los parametros y captura mayormente las dependencias de corto plazo

Como resolver el problema:

- Activation function
- Inicializacion de pesos
- Nuevas arquitecturas

Vanishing gradient problem

Multiplica muchos numeros pequeños




Errores debido a que para aquellos pasos lejanos en el tiempo, los gradientes se hacen cada vez mas pequeños



Sesga los parametros y capture mayormente las dependencias de corto plazo

- Activation function



Vanishing gradient problem

Multiplica **muchos numeros pequeños**



Errores debido a que para aquellos pasos lejanos en el tiempo, los gradientes se hacen cada vez mas pequeños



Sesga los parametros y capture mayormente las dependencias de corto plazo

Inicialización de pesos

$$I_n = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{pmatrix}$$

Iniciarizar **weights** con la matriz identidad
Iniciarizar **biases** com cero

Previene que los gradientes se vayan a cero

Vanishing gradient problem

Multiplica muchos numeros pequenos




Errores debido a que para aquellos pasos lejanos en el tiempo, los gradientes se hacen cada vez mas pequenos



Sesga los parametros y capture mayormente las dependencias de corto plazo

- Nuecas arquitecturas

Solucion mas robusta




Usa unidades recurrentes con compuertas que son mas complejas para controlar el flujo de la informacion que se esta pasando a lo largo de la secuencia

Las compuertas agregan o remueven informacion selectivamente dentro de cada unidad recurrente

Long short term memory (LSTM)

Gates: 1) Forget 2) Input (store) 3)Update 4) Output




Como trabaja:

- 1) Mantiene una celda de estado
- 2) Usa compuertas (gates) para controlar el flujo de la información.
 - Forget gate: se deshace de la información poco relevante
 - Almacén información relevante de la entrada actual (store/input gate)
 - Selectivamente hace un update de la celda de estado
- Output gate: retorna la versión filtrada de la celda de estado
- 3) Backpropagation TT con flujo de gradientes parcialmente sin interrupciones

Long short term memory (LSTM)

Gates: 1) Forget 2) Input (store) 3)Update 4) Output



How it works:

- 1) Maintain a **cell state**
- 2) Use gates to control the flow of information
 - **Forget gate** gets rid of irrelevant information
 - Store relevant information from the current input
 - Selectively **update cell** state
 - **Output gate** returns a filtered version of the cell state
- 3) Backpropagation TT with partially uninterrupted gradient flow

LSTM "no todo lo que brilla es oro"

- << vanishing gradient problem pero no lo elimina completamente.
- >> Costo computacional
- Es afectado por diferentes inicializaciones aleatorias de los pesos
- Drop-out es difícil de implementar
- Tiende a hacer overfitting
- << performance en problemas con dependencias temporales muy largas



Convolutional Neural Networks

Received July 15, 2019; accepted July 30, 2019; date of publication August 5, 2019; date of current version August 19, 2019.
Digital Object Identifier 10.1109/ACCESS.2019.2933060

Automated Individual Pig Localisation, Tracking and Behaviour Metric Extraction Using Deep Learning

JAKE COWTON^{1,2}, ILIAS KYRIAZAKIS², AND JAUME BACARDIT²
¹School of Computing, Newcastle University, Newcastle upon Tyne NE1 7RU, U.K.
²School of Natural and Environmental Sciences, Agriculture, Newcastle University, Newcastle upon Tyne NE1 7RU, U.K.
Corresponding author: Jake Cowton (j.cowton2@newcastle.ac.uk)

This work was supported by the European Commission under the European Union Framework Programme for Research and Innovation Horizon 2020 under Grant 633531. The work of J. Bacardit was supported by the Engineering and Physical Science Research Council under Grant EP/N031962/1 and Grant EP/M020576/1.





FIGURE 8. Four sample images from our pig detection test set processed by the Faster R-CNN with the feature extraction layers pre-trained on ImageNet, the rest pre-trained on Pascal Visual Object Classes Challenge 2007 and an additional fully-connected layer for the pig dataset. Detections to the left of the red wall are ignored. The top left image is from the low-light test segment. The top right image is from the densely packed test-segment. The bottom left image is from the overexposed test segment. The bottom right image is from the many pigs test segment.


CNNs

Computer Vision.

- Facial detection and recognition
- Healthcare, medicine and biology
- Self-driving vehicles



Convolutional Neural Networks




What the computer sees																
157	153	174	169	150	152	129	151	172	161	155	156	155	182	163	74	75
155	182	163	74	75	62	33	17	110	210	180	154	180	180	50	14	34
180	180	50	14	34	6	16	53	48	105	159	181	206	109	6	124	131
206	109	6	124	131	111	120	204	166	15	56	180	194	68	137	261	297
194	68	137	261	297	289	299	228	227	87	71	201	172	105	207	233	233
172	105	207	233	233	214	220	239	228	98	74	206	188	88	179	209	185
188	88	179	209	185	215	211	158	193	75	20	169	189	97	165	84	10
189	97	165	84	10	168	134	11	31	62	22	148	199	168	191	193	158
199	168	191	193	158	227	178	143	182	105	36	190	205	174	155	252	236
205	174	155	252	236	231	149	178	228	43	95	234	190	216	116	149	236
190	216	116	149	236	187	85	150	79	38	218	241	190	224	147	108	227
190	224	147	108	227	216	127	103	56	101	255	234	190	214	173	66	103
190	214	173	66	103	143	96	60	2	109	249	215	187	196	235	75	1
187	196	235	75	1	81	47	0	6	217	255	211	183	202	237	145	0
183	202	237	145	0	9	12	108	200	138	243	236	195	206	123	207	177
195	206	123	207	177	121	123	200	175	13	96	218					

An image is just a matrix of numbers [0,255]!
i.e., 1080x1080x3 for an RGB image

Features

(Como en otros problemas con NN)



- Regresion
- Clasificacion
- Deteccion de features Ede alto nivel



Extraer las features a partir de los
datos!
Aprender jerarquias de features!

Fully connected NN

- Input:
- 2D image
 - vector de valores de pixeles
(flatten the image)



Fully connected:



- Conectar una neurona de una capa oculta a todas las neuronas en la entrada
- No mantiene informacion espacial!
- Muchos parametros

Como le agregamos una estructura espacial a la entrada?

Using spatial structure

Input:


- 2D image
- Arreglo de pixeles



Sliding window para definir las conexiones (conecta un parche en la capa de entrada a una sola neurona)

La clave: como podemos **pesar** cada parche para detectar una feature particular?

Using spatial structure





- Filtro de size 4x4: son 16 pesos diferentes
- Aplicar el mismo filtro a 4x4 parches en la entrada
- Mover 2 pixels para el proximo parche




Esta operacion se llama **convolucion**

- Aplicar un set de pesos (un filtro) para extraer features locales
- Usar multiples filtros para extraer diferentes features
- Cada filtro **comparte espacialmente parametros**

Convolutional operation




Convolutional operation



Convolutional neural networks (CNN)




- for classification



- **Convolucion:** aplicar filtros para generar 'feature maps'
- **Non-linealidad:** mas usado ReLU
- **Pooling:** operacion de submuestreo (downsampling) en cada mapa de features

**Train model with image data.
Learn weights of filters in convolutional layers.**

Convolutional layer





1. Aplicar una ventana de pesos

2. calcular las combinaciones lineales


Para cada neurona de la capa oculta:

- Toma las entradas del parche
- Calcula la suma ponderada
- Aplica bias

3. se activa con una función no lineal



Pooling




max pool con filtro 2x2
Y un paso de 2



6	8
3	4


- Reduce la dimensionalidad
- Invarianza espacial

CNN: classification example




- Hasta este punto, la ultima salida es el i -esimo feature map

CNN: classification example



$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$



logits = predicciones no normalizados (aun)
del modelo

CONV and POOL layers putput **high-level features**
input

Salida de las capas CONV y POOL >> **features de alto nivel**


Una capa fully connected usa estas **features** para **clasificar** la image

Expresa la **salida como una probabilidad** (de la imagen a una cierta clase)



¿Y AHORA QUÉ?

Deployment




Hidden Technical Debt in Machine Learning Systems, D. Sculley et.al (2015)

- Tiene todos los problemas de deployment y mantenimiento de un código tradicional más los problemas específicos de ML
- El modelo es solo una pequeña parte de un sistema operacional


- How easily can an entirely new algorithmic approach be tested at full scale?
- What is the transitive closure of all data dependencies?
- How precisely can the impact of a new change to the system be measured?
- Does improving one model or signal degrade others?
- How quickly can new members of the team be brought up to speed?

Robustness & Trusworthiness

Barredo +, 2019




Robustness & Trusworthiness




Robustness & Trusworthiness

Class Imbalance




Mitigating Class Imbalance





Robustness & Trusworthiness

Latent features



Robustness & Trusworthiness

Uncertainty




```
num_ensembles = 5
for i in range(num_ensembles):
    model = create_model(...)
    model.fit(...)

raw_predictions = [models[i].predict(x)
                   for i in range(num_ensembles)]
mu = np.mean(raw_predictions)
uncertainty = np.var(raw_predictions)
```



An application


- 2 stages: a) **single station forecasting (ML)**;
b) extended forecasting
- 3 meridional sectors covering low, mid & high latitude
- Covering land & oceanic regions
- **Input: TEC from GIMs + External input (Kp)**



Objectives:

- Global TEC forecast 24 hs ahead using DL
- Propose a semi-operative prototype

Molina +(submitted)




Cesaroni +, 2020



An application


- 2 stages: a) **single station forecasting (ML)**;
b) extended forecasting
- 3 meridional sectors covering low, mid & high latitude
- Covering land & oceanic regions
- **Input: TEC from GIMs + External input (Kp)**



Objectives:

- Global TEC forecast 24 hs ahead using DL
- Propose a semi-operative prototype

Molina + (under review)



Cesaroni +, 2020




Data preparation & Feature selection


- Dataset:
 - 2005 - 2016
 - splitting strategy: 99% (99 train/1val) - 1% test (~43 days)
 - + cases study: geomagnetic storms in 2017
- Resolution (re-sampling):
 - TEC from GIMs - 2 hs resolution
 - Kp - 3hs resolution > K Nearest-neighbor interpolation
- Smart weight initialization (kernel initialization): GlorotNormal distribution + proper activation function (e.g. tanh).

Loosely physics-informed approach

St 01 TEC - dataset (2005 - 2016)




TEC - single ST Histogram






ML modelling

- 3 ML techniques:
 - 2 RNNs (LSTM & GRU)
 - CNN (1D)
- Time series
- Hyperparameter tuning:
grid search



- # hidden layers (5,10,15,20,50,100 cells)
- batch size (16,32,64,128)
- #epochs (iterations) (5,10,15,20,30,40,50,100,200, 500,1000)




RNNs:

- Maintain order
- Memory (h_t)
- Backpropagation through time
- Prone to overfitting, vanishing gradient problem
- LSTM & GRU -> gated cells -> long-term but not that long

CNN:



- kernel size = 2





ML modelling

- Why these results?
 - LSTM & GRU -> difficult to catch fast changes and peaks
 - CNN (1D) -> spatial relationship = short term relationships




- Forecasting 24 hs ahead (quiet day)
- RMSE < 3 TECu
- CNN best at any station (- St16,17,18 -> TECu<=1 -> quiet day)
- Low lat + oceanic stations -> + challenging




ML modelling

- In general: in SWx, few extreme cases (unbalanced datasets) → forecasting may fail when new data arrives (generalization is a problem) → Incremental learning

RMSE



Test set -> 43 days with the basic models



Test set -> 43 days with the models + incremental learning (updating each 24 hs)



ML modelling

- We considered cases study from 2017 under different geomagnetic conditions

$$\text{Global } \Delta TEC = \frac{1}{st} \sum \Delta TEC$$

