



Ciencia de Datos

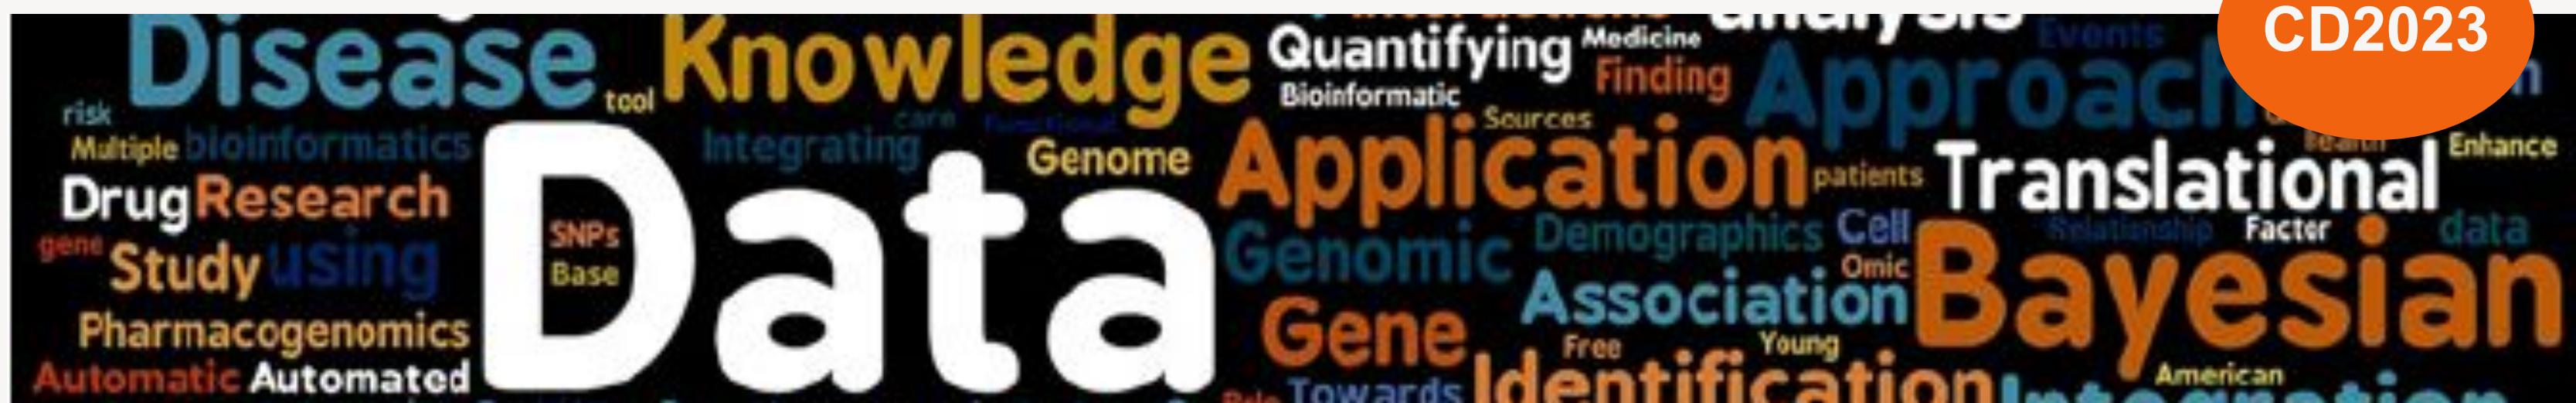
2025

Introducción a la Ciencia de
Datos

OPTATIVA - LICENCIATURA EN INFORMÁTICA
FACET-UNT

Ciencia de Datos: Fundamentos y
Herramientas

CURSO DE POSTGRADO - FACET-UNT



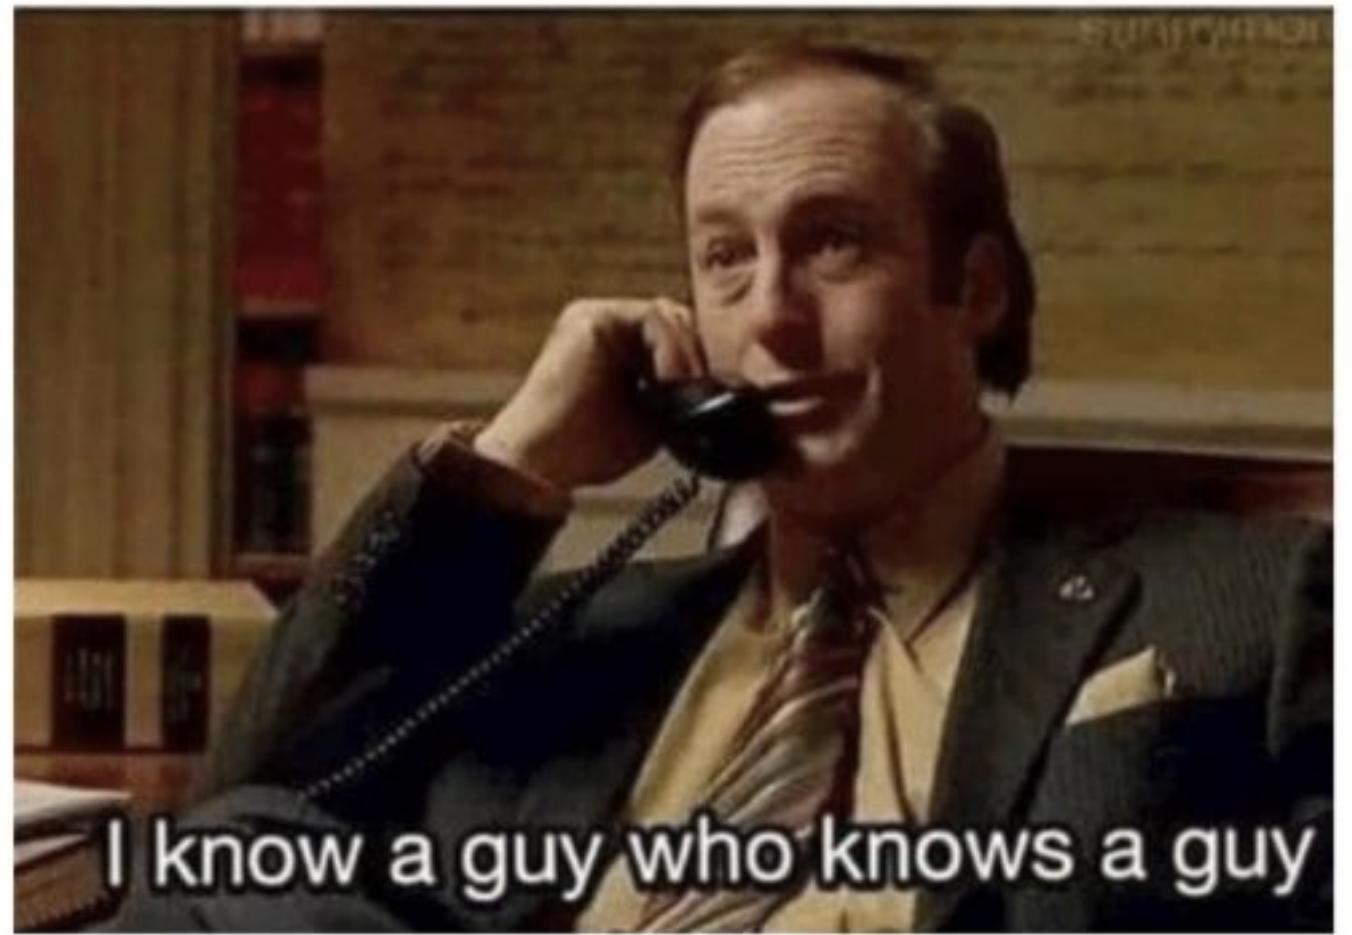
Artificial Neural Networks

ANN

NN

How Neural Networks work?

Neurons:



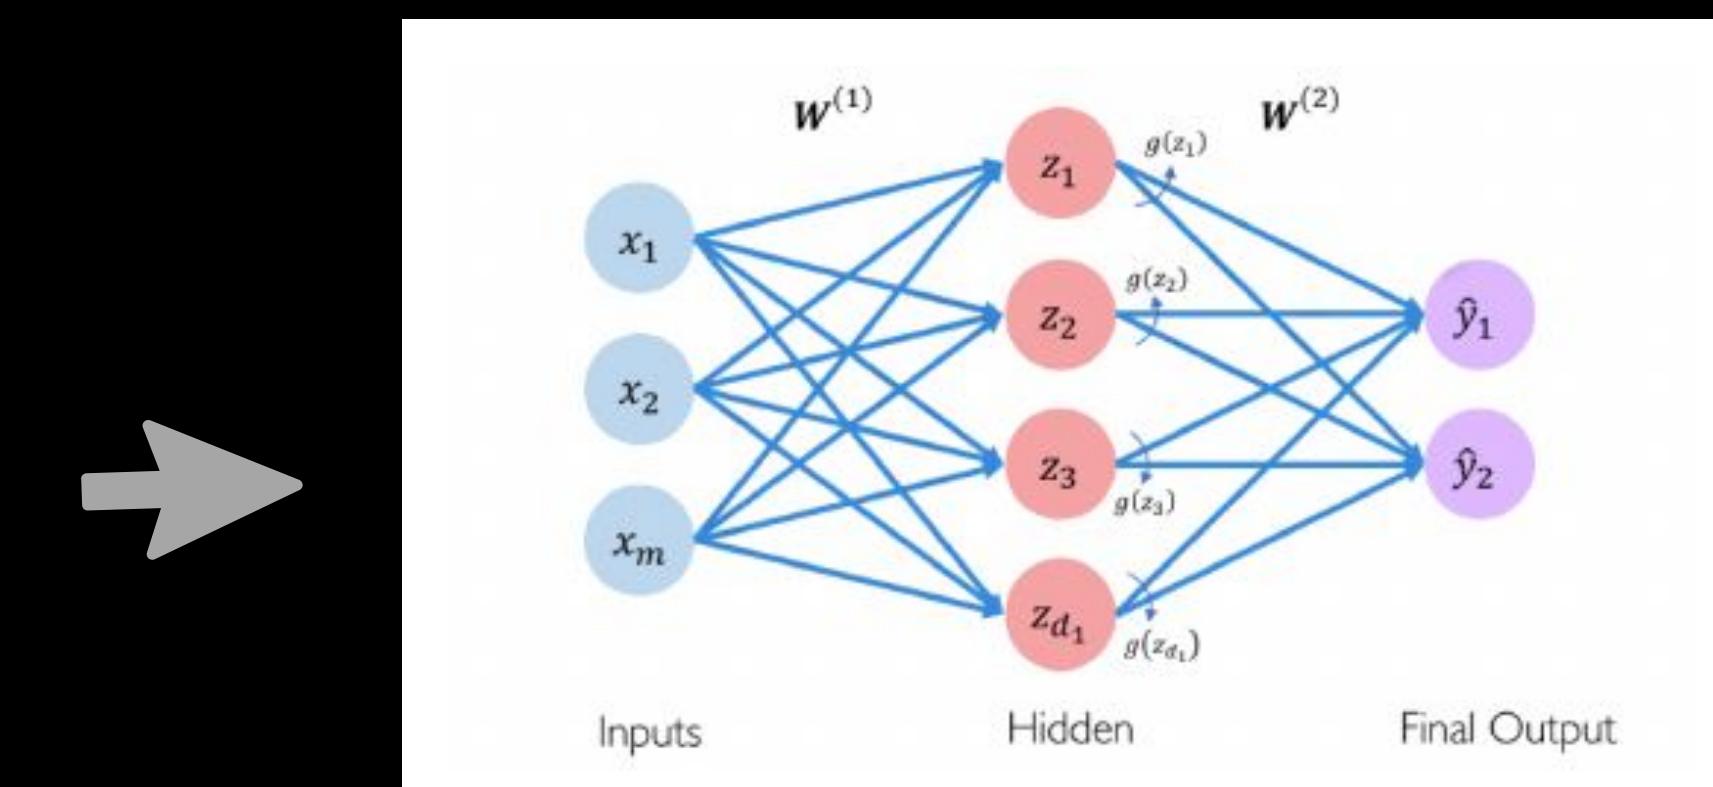
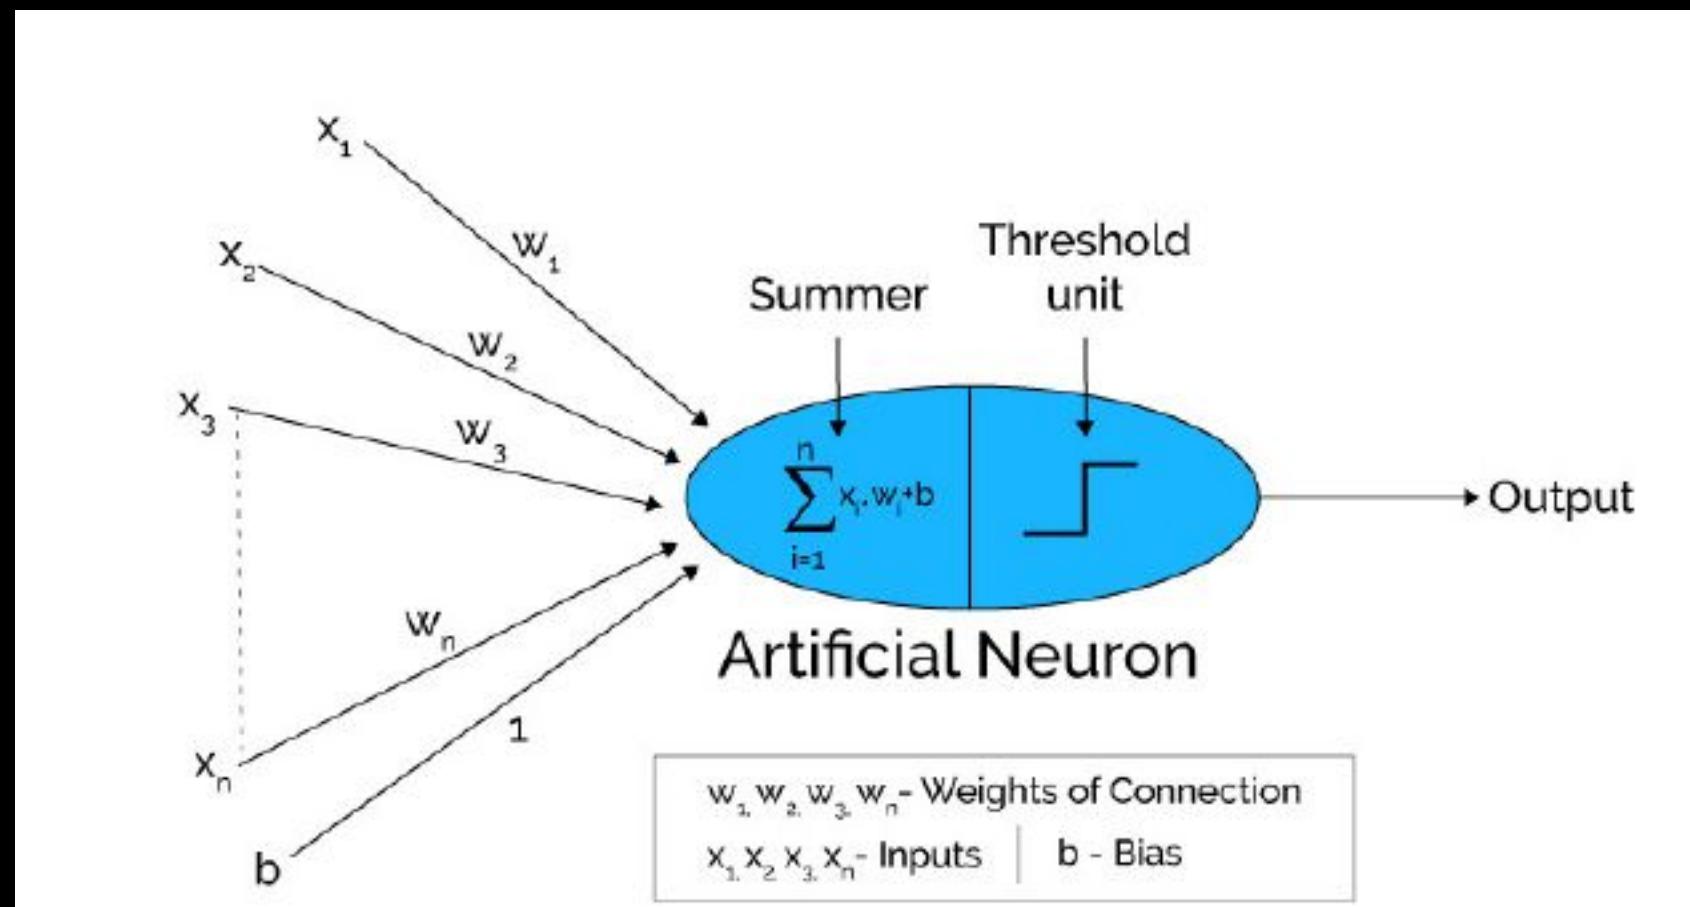
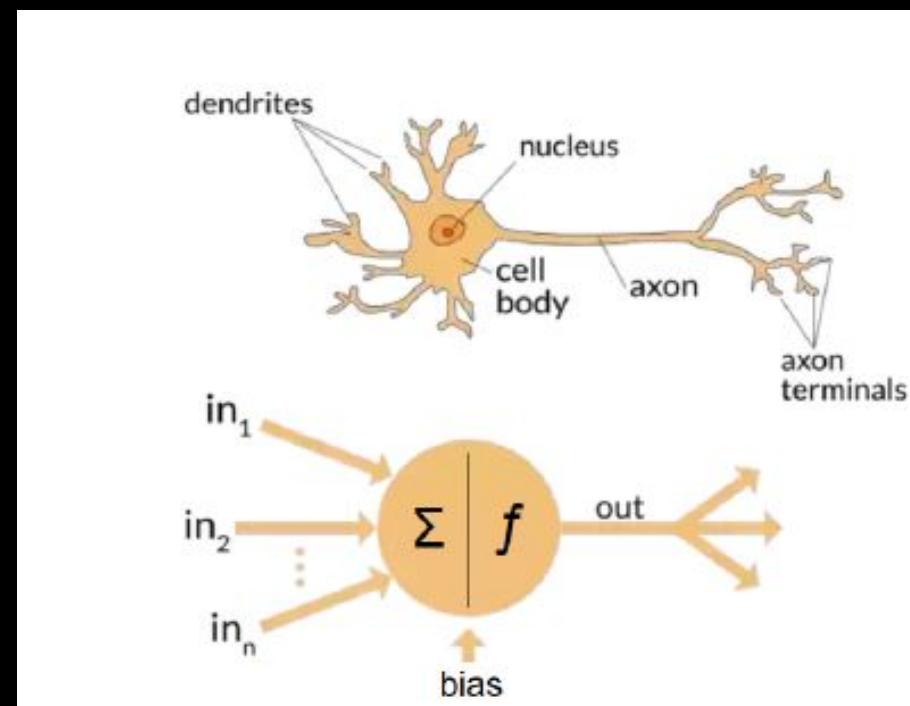
I know a guy who knows a guy

Redes Neuronales Artificiales

RNA

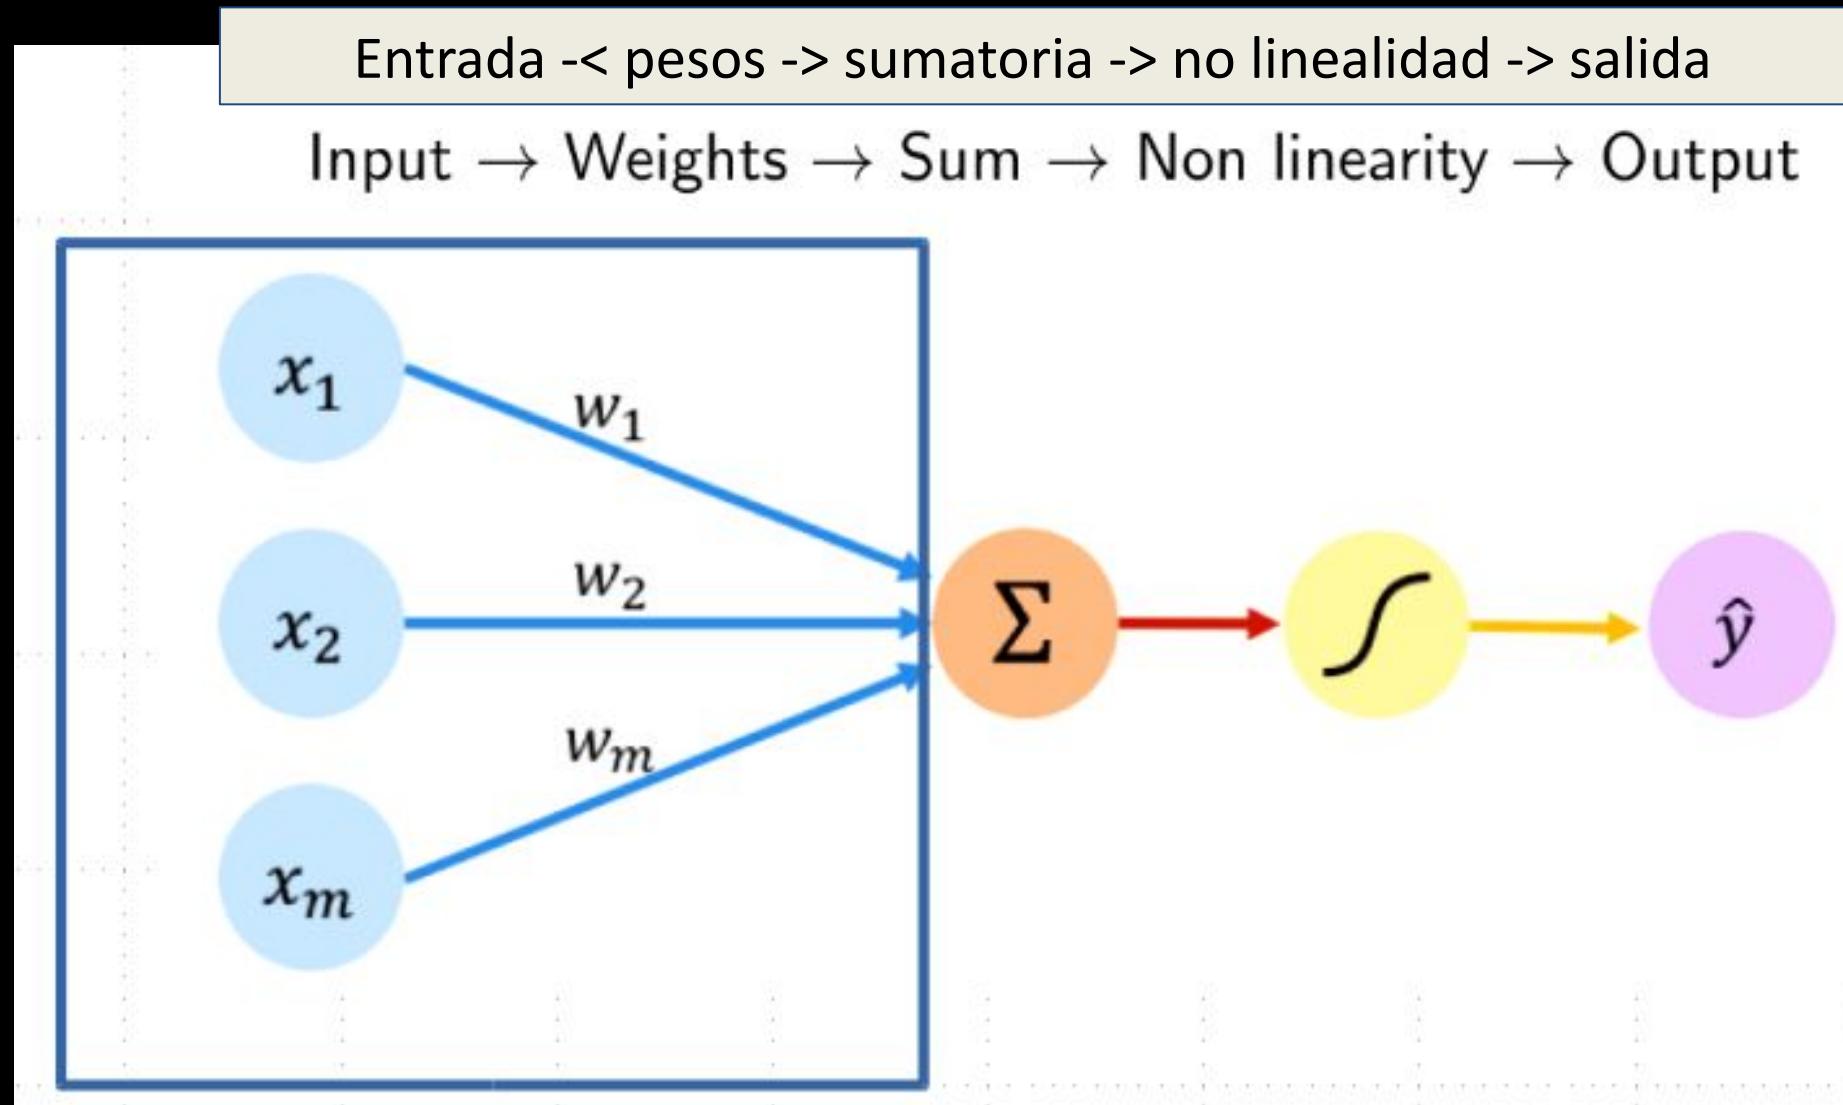
ANN

- Data-driven modeling / modelado basado en datos > pobre capacidad de generalización
- Inspirado en las redes neuronales del cerebro
- Para resolver un gran número de problemas muy complejos: reconocimiento facial, reconocimiento de la escritura, pronóstico del clima, etc.
- Black-box modelling (?) > basado en la composición de neuronas interconectadas > difícil entender las predicciones



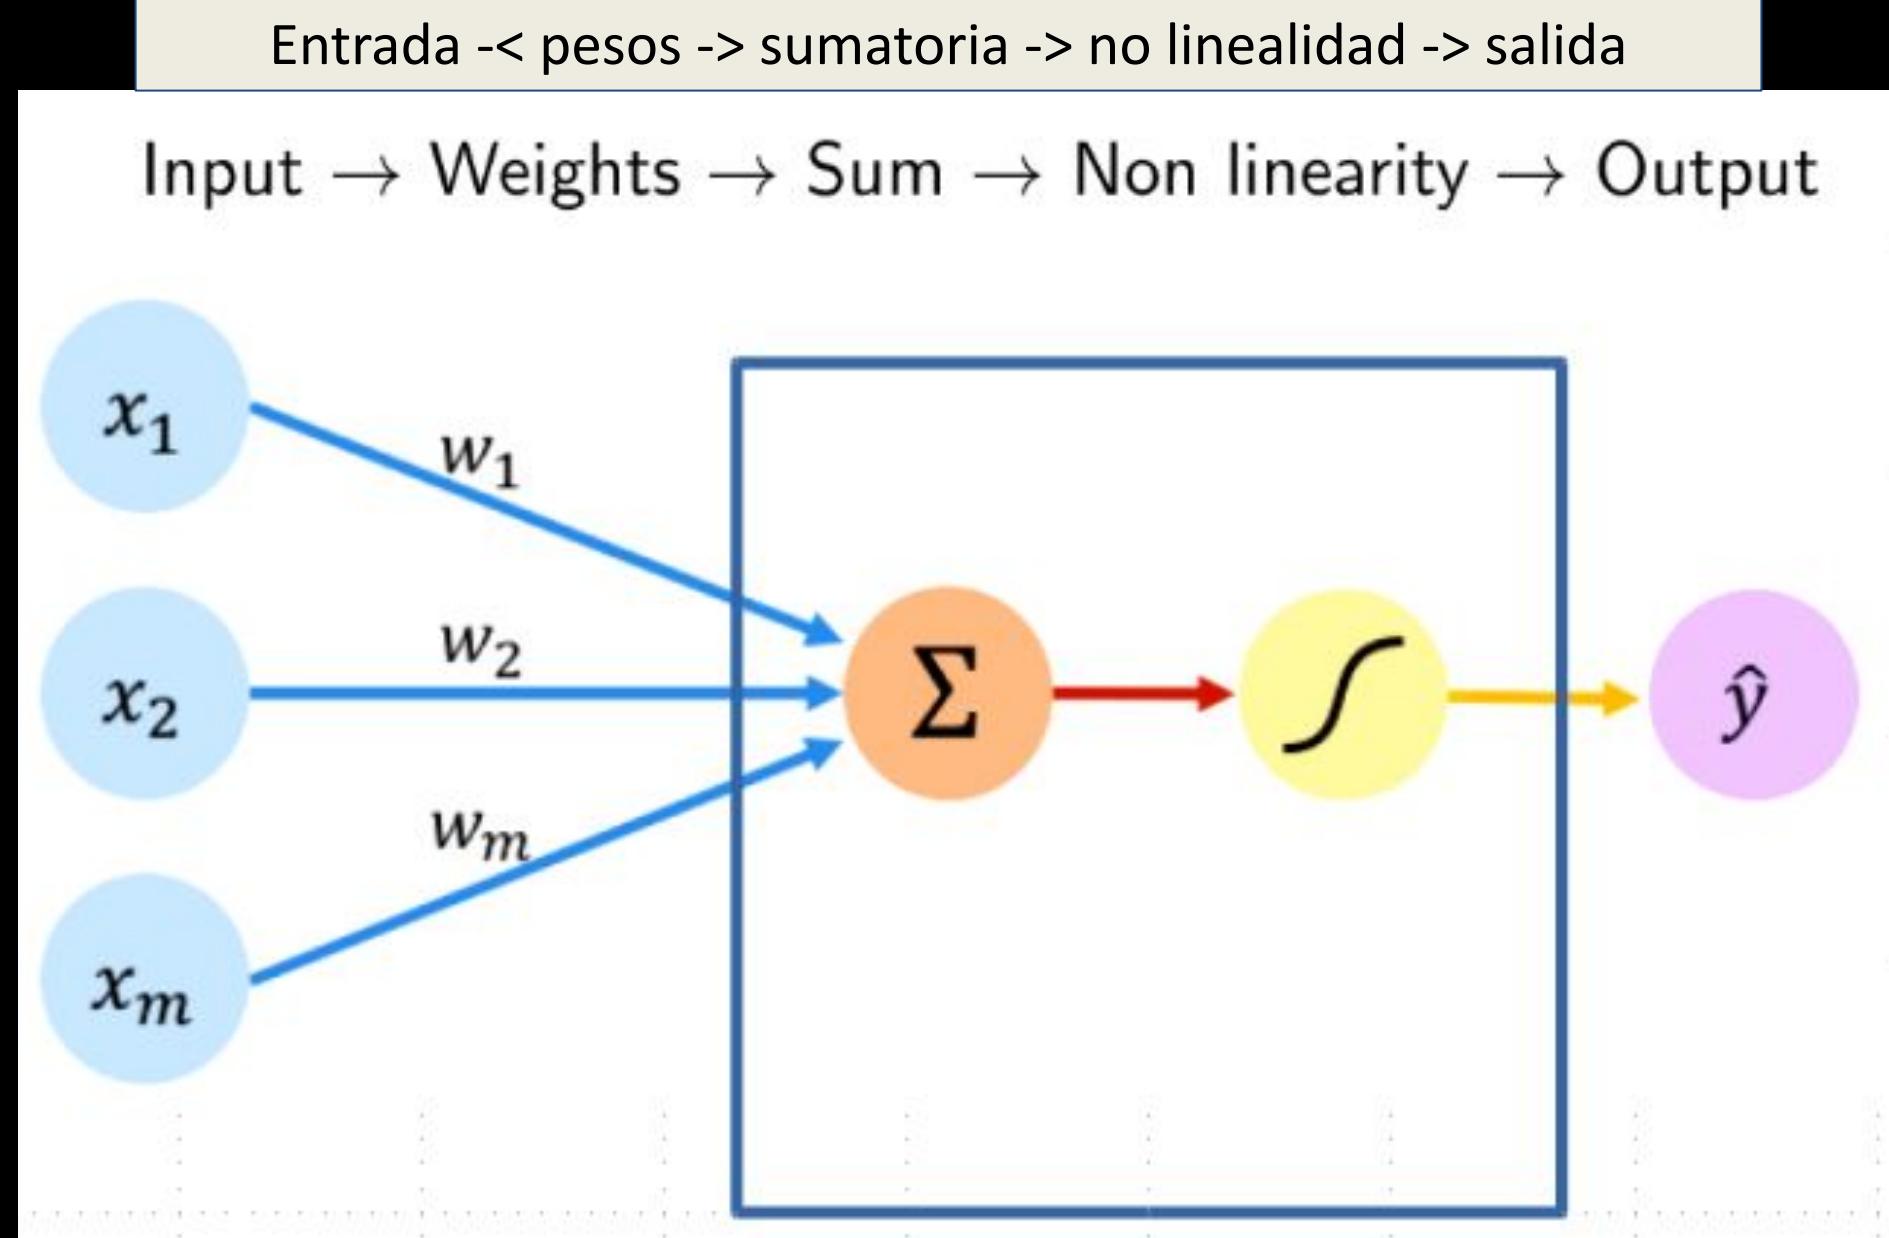
Perceptron

- Unidad fundamental de una red neuronal



- cada entrada = una ‘feature’
- se aplica un peso(w) a cada entrada (x_i)
- los pesos son inicializados y actualizados de manera aleatoria durante el entrenamiento

Perceptron



Output

Linear combination of inputs

$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function

Bias

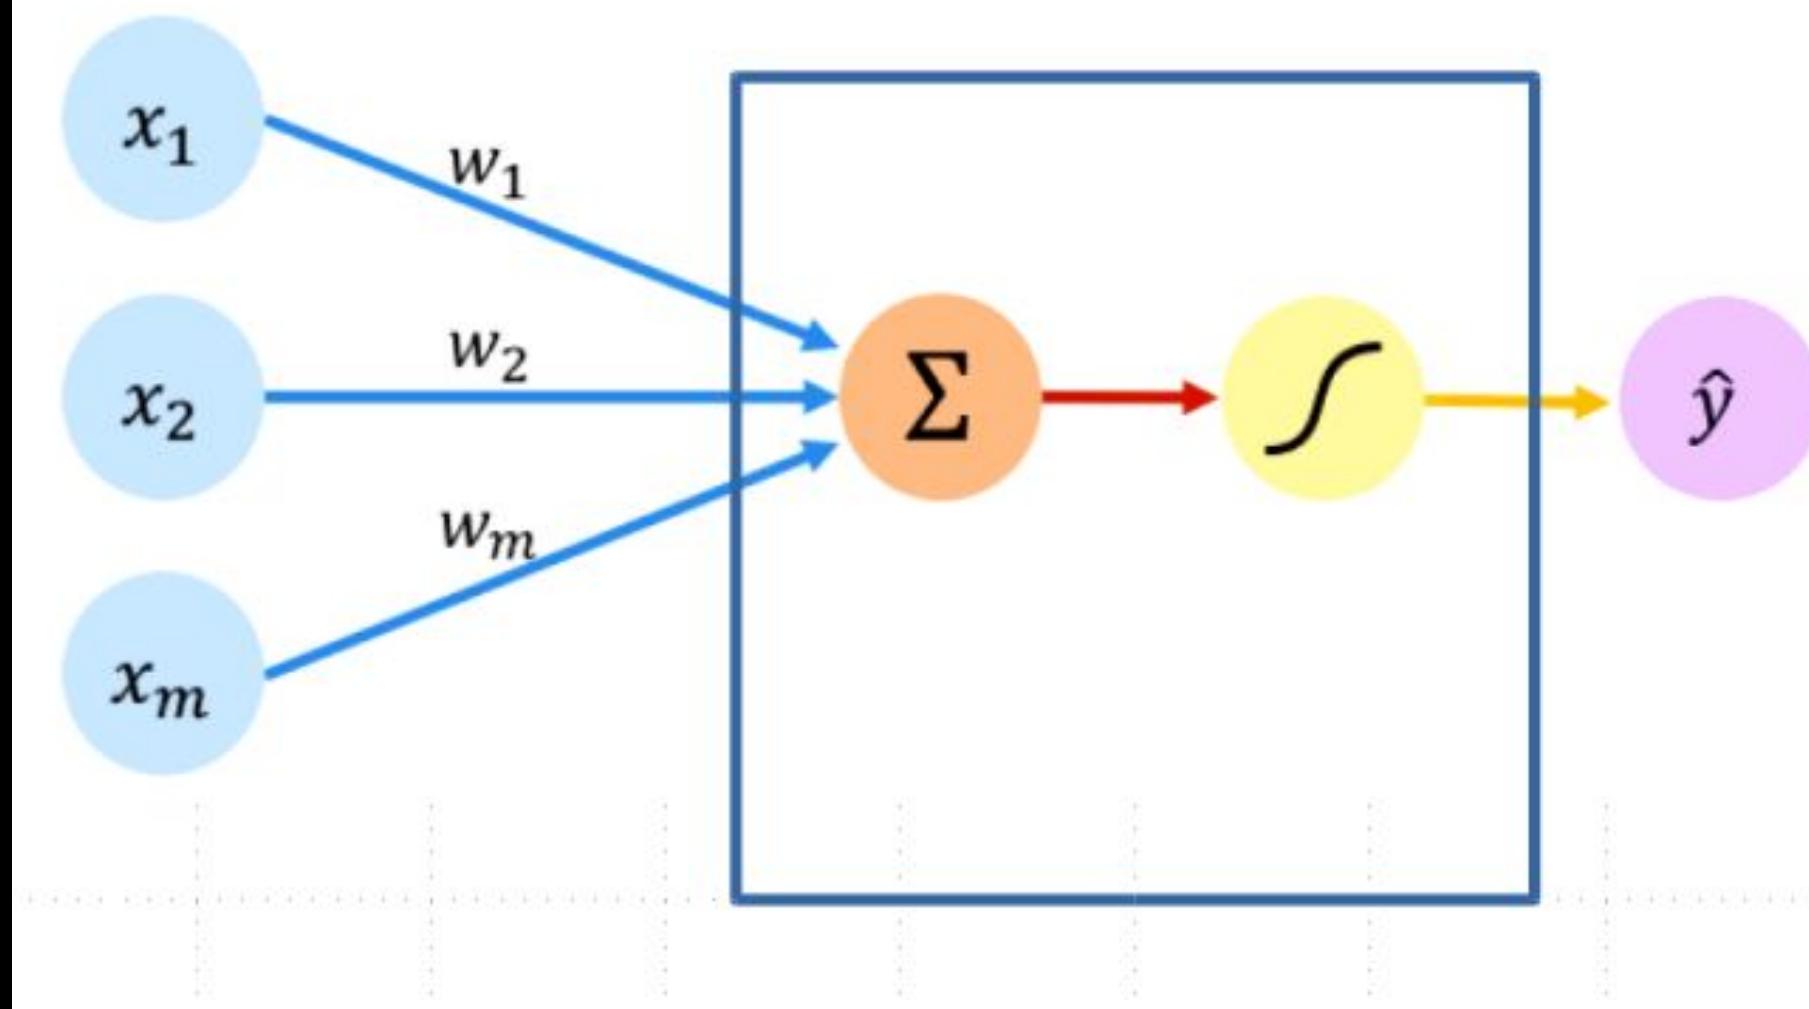
The diagram shows the mathematical formula for a perceptron's output. The output \hat{y} is the result of applying a non-linear activation function g to the linear combination of inputs. The linear combination is calculated as $w_0 + \sum_{i=1}^m x_i w_i$. A purple arrow points down to the output \hat{y} . A red arrow points down to the term $w_0 + \sum_{i=1}^m x_i w_i$ and is labeled "Linear combination of inputs". A green arrow points to the term w_0 and is labeled "Bias". A yellow arrow points to the activation function g and is labeled "Non-linear activation function".

- capa oculta o 'hidden layer'
- Bias/sesgo: término de ajuste (muchas veces seteado en 1)
- la función de activación es aplicada a la combinación lineal de las entradas pesadas (+ el bias)

Perceptron

Entrada -< pesos -> sumatoria -> no linealidad -> salida

Input → Weights → Sum → Non linearity → Output



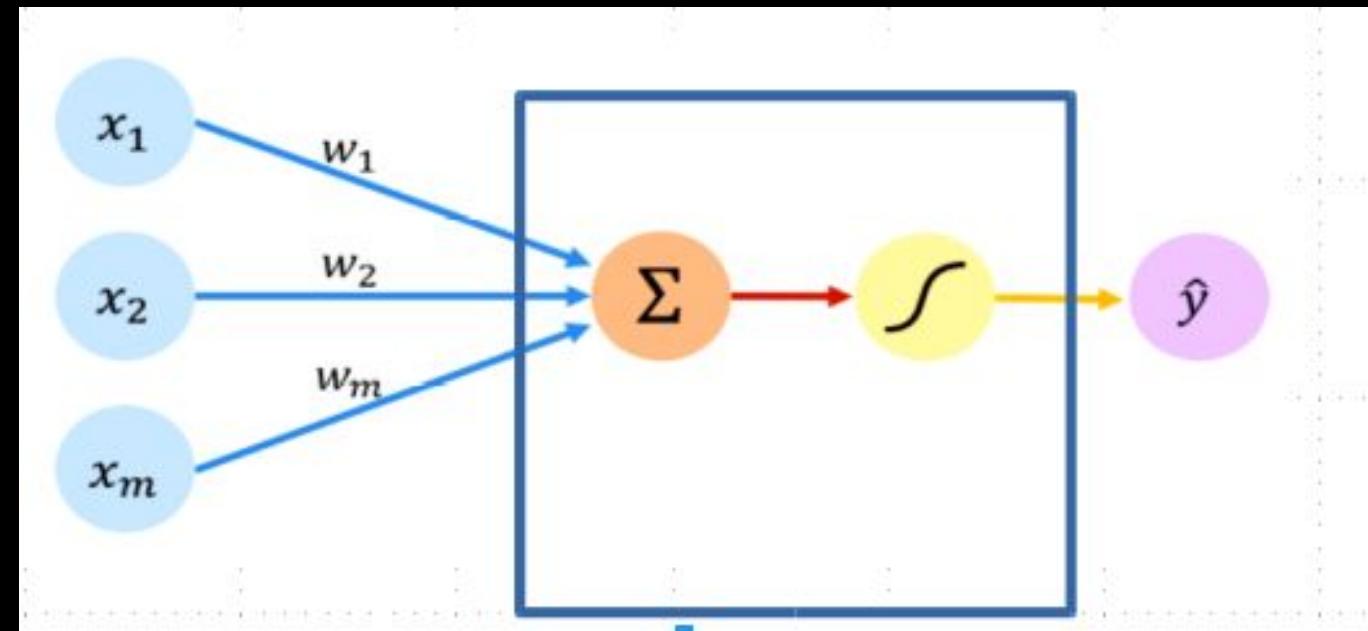
$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

$$\hat{y} = g (w_0 + \mathbf{X}^T \mathbf{W})$$

where: $\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$ and $\mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$

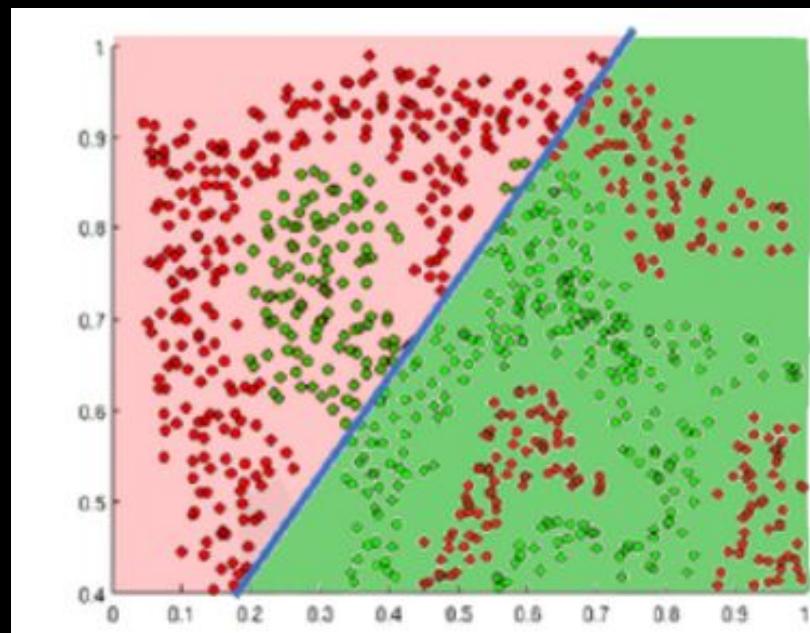
Función de activación

- Provee de la capacidad de agregar no-linealidad al modelo

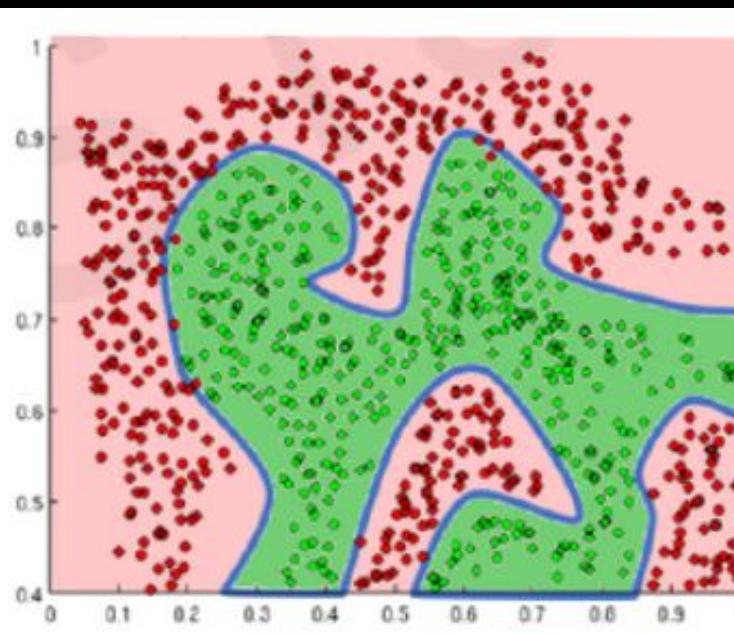


$$\hat{y} = g(w_0 + X^T W)$$

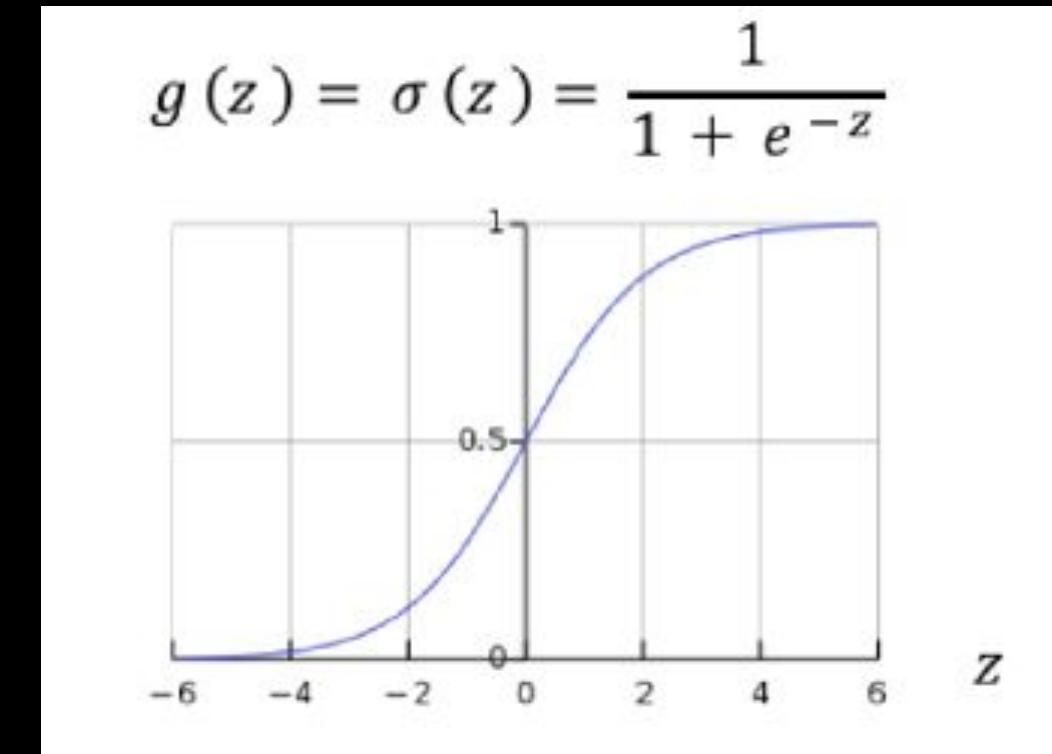
- Una de las más usadas: función sigmoide que produce valores entre 0 y 1



Linear activation functions produce linear decisions no matter the network size



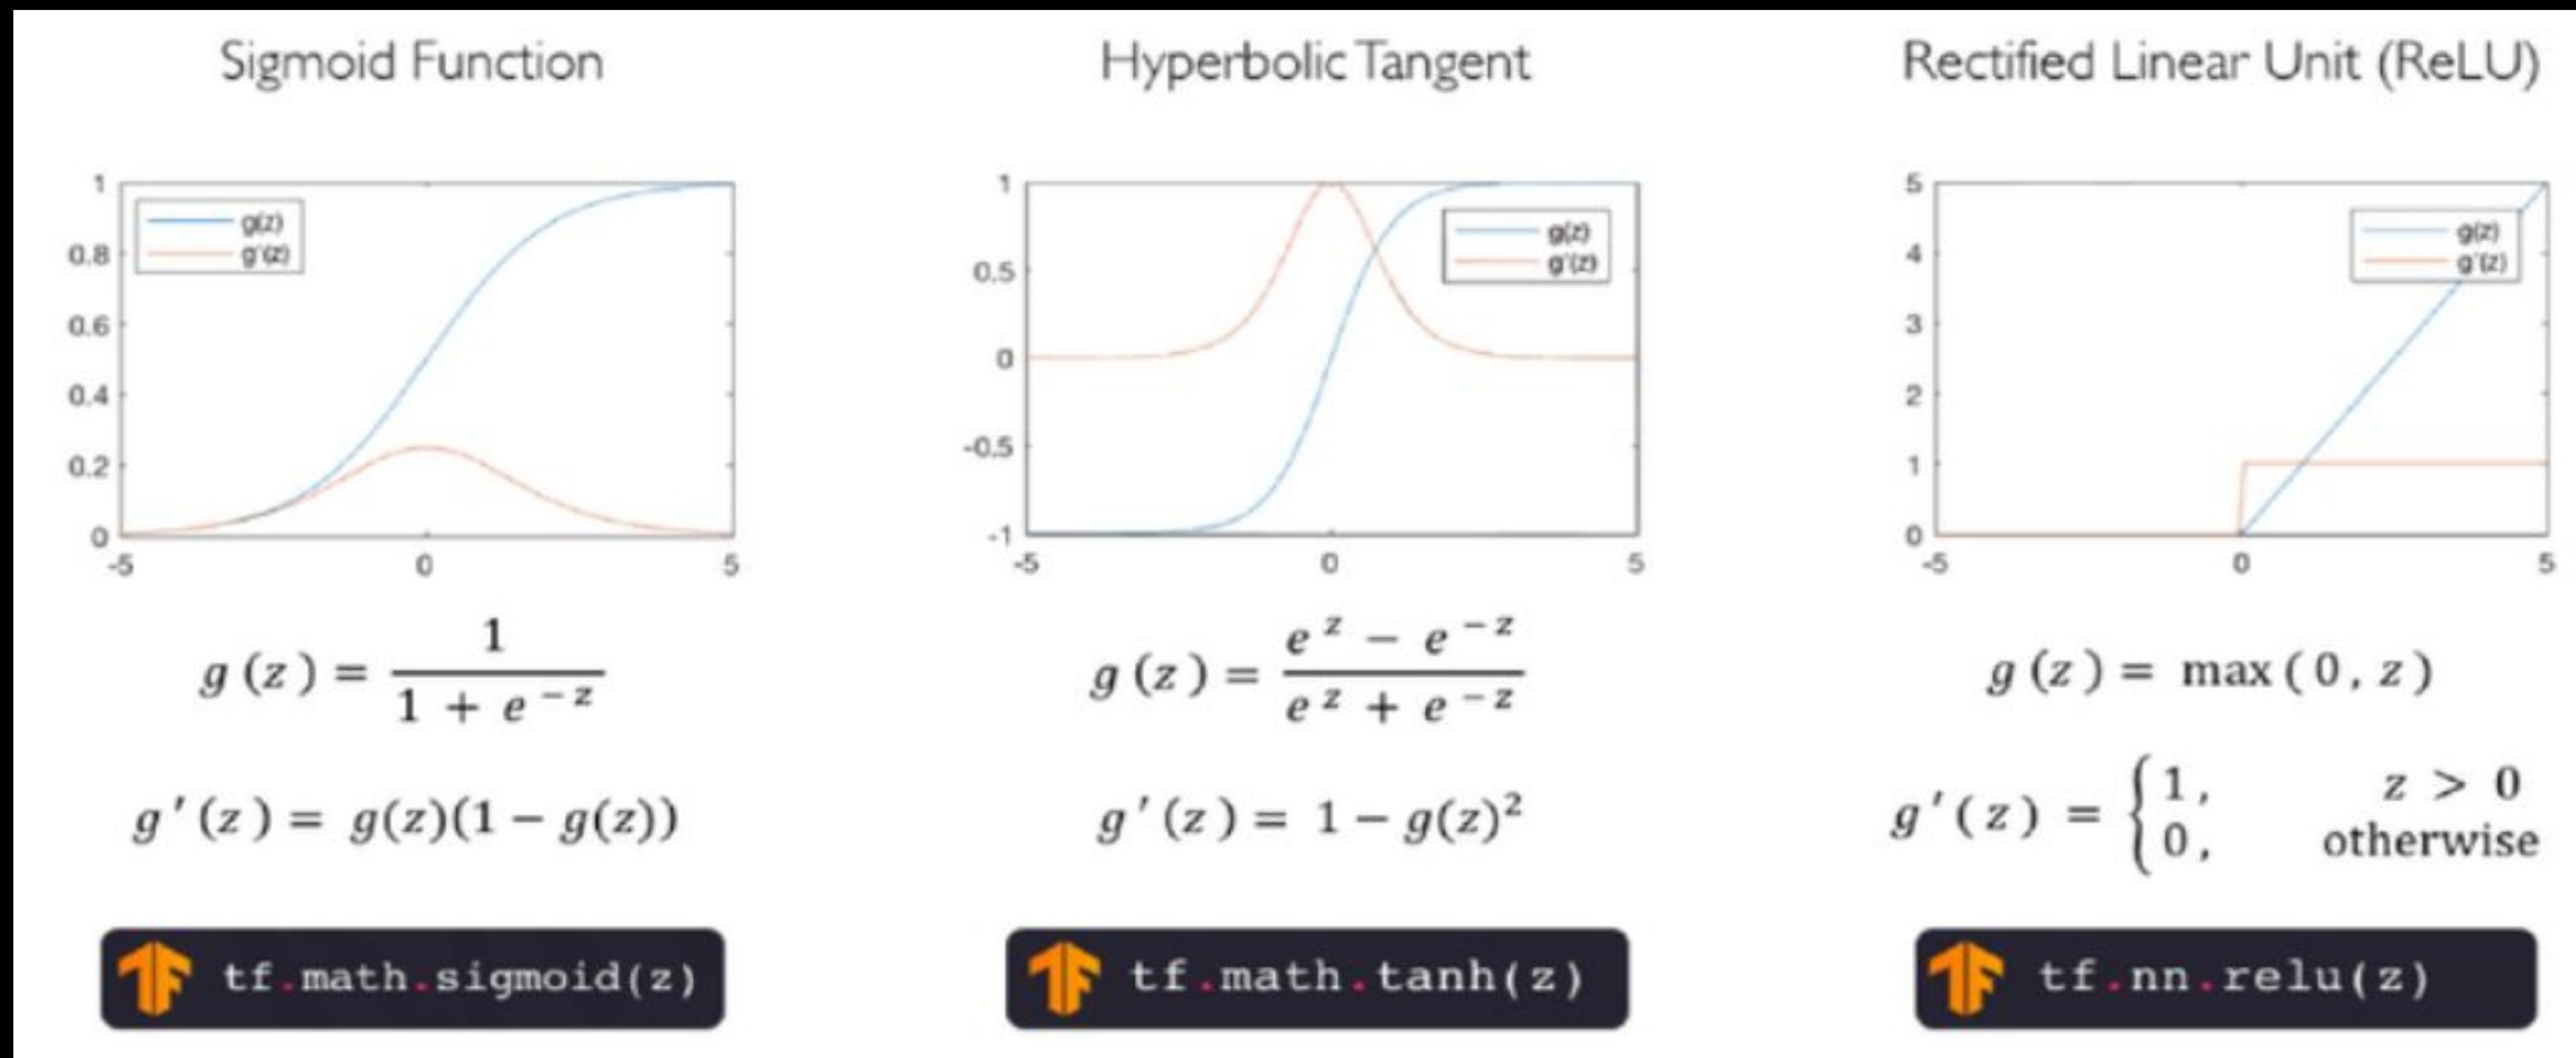
Non-linearities allow us to approximate arbitrarily complex functions



Funcion de activacion

- vamos a ver que no solo la función es importante sino también su derivada
- Es una configuración (hiperparametro) que se debe seleccionar ante del entrenamiento.

$$\hat{y} = g(w_0 + X^T W)$$



$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$

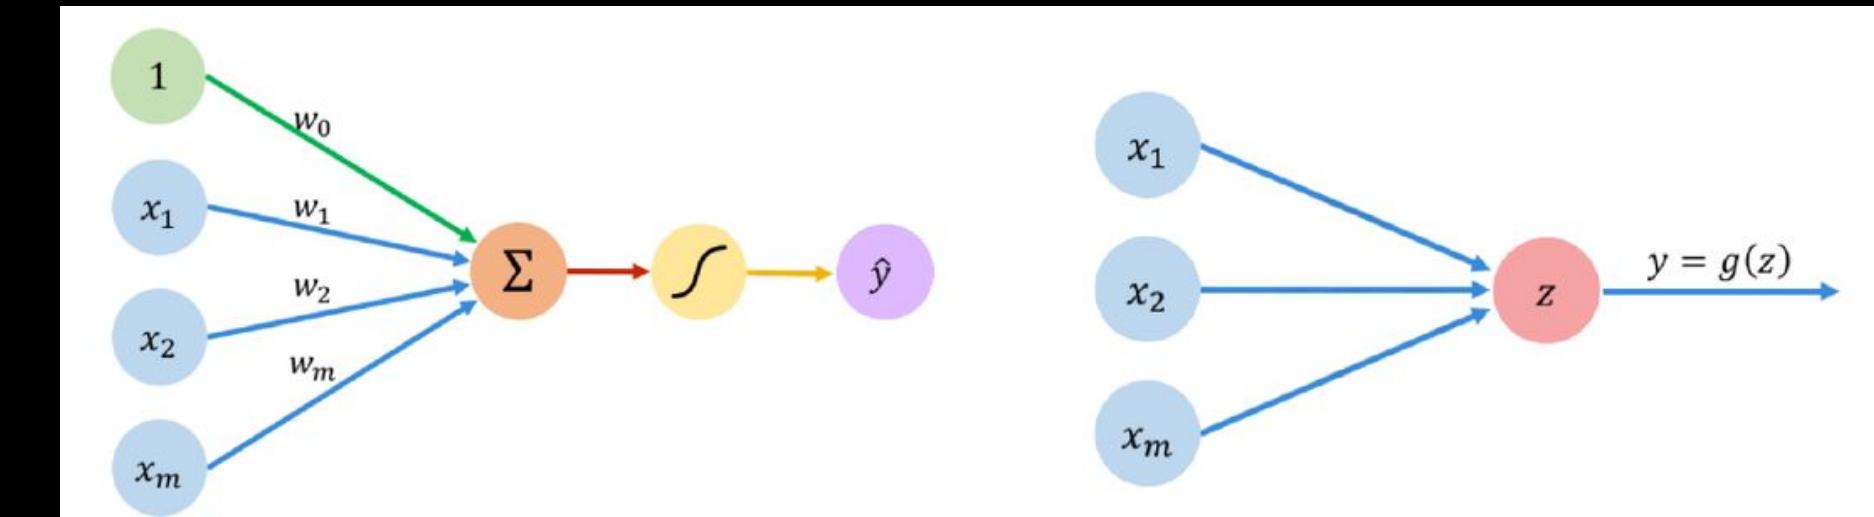
- Elegir una función de activación acorde al problema y a los datos es crucial para que el modelo tenga una buena performance

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

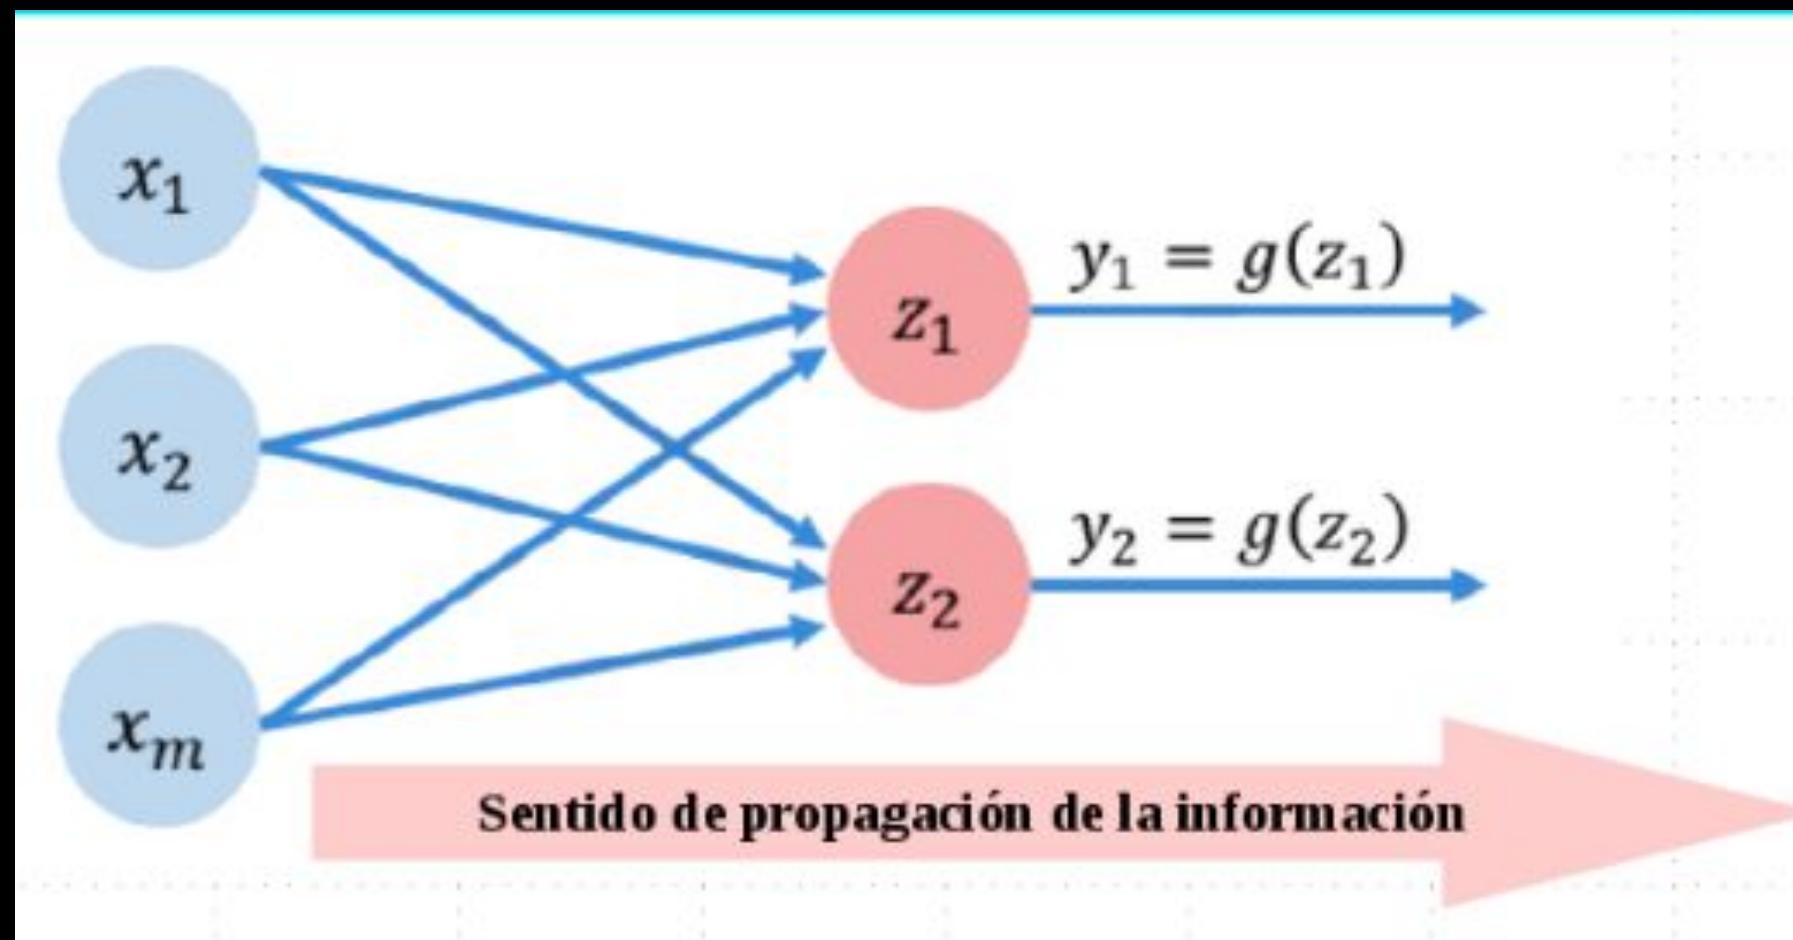
Feed Forward (FF)

- Perceptron simple
(n inputs → 1 output)

$$z = w_0 + \sum_{j=1}^m x_j w_j$$



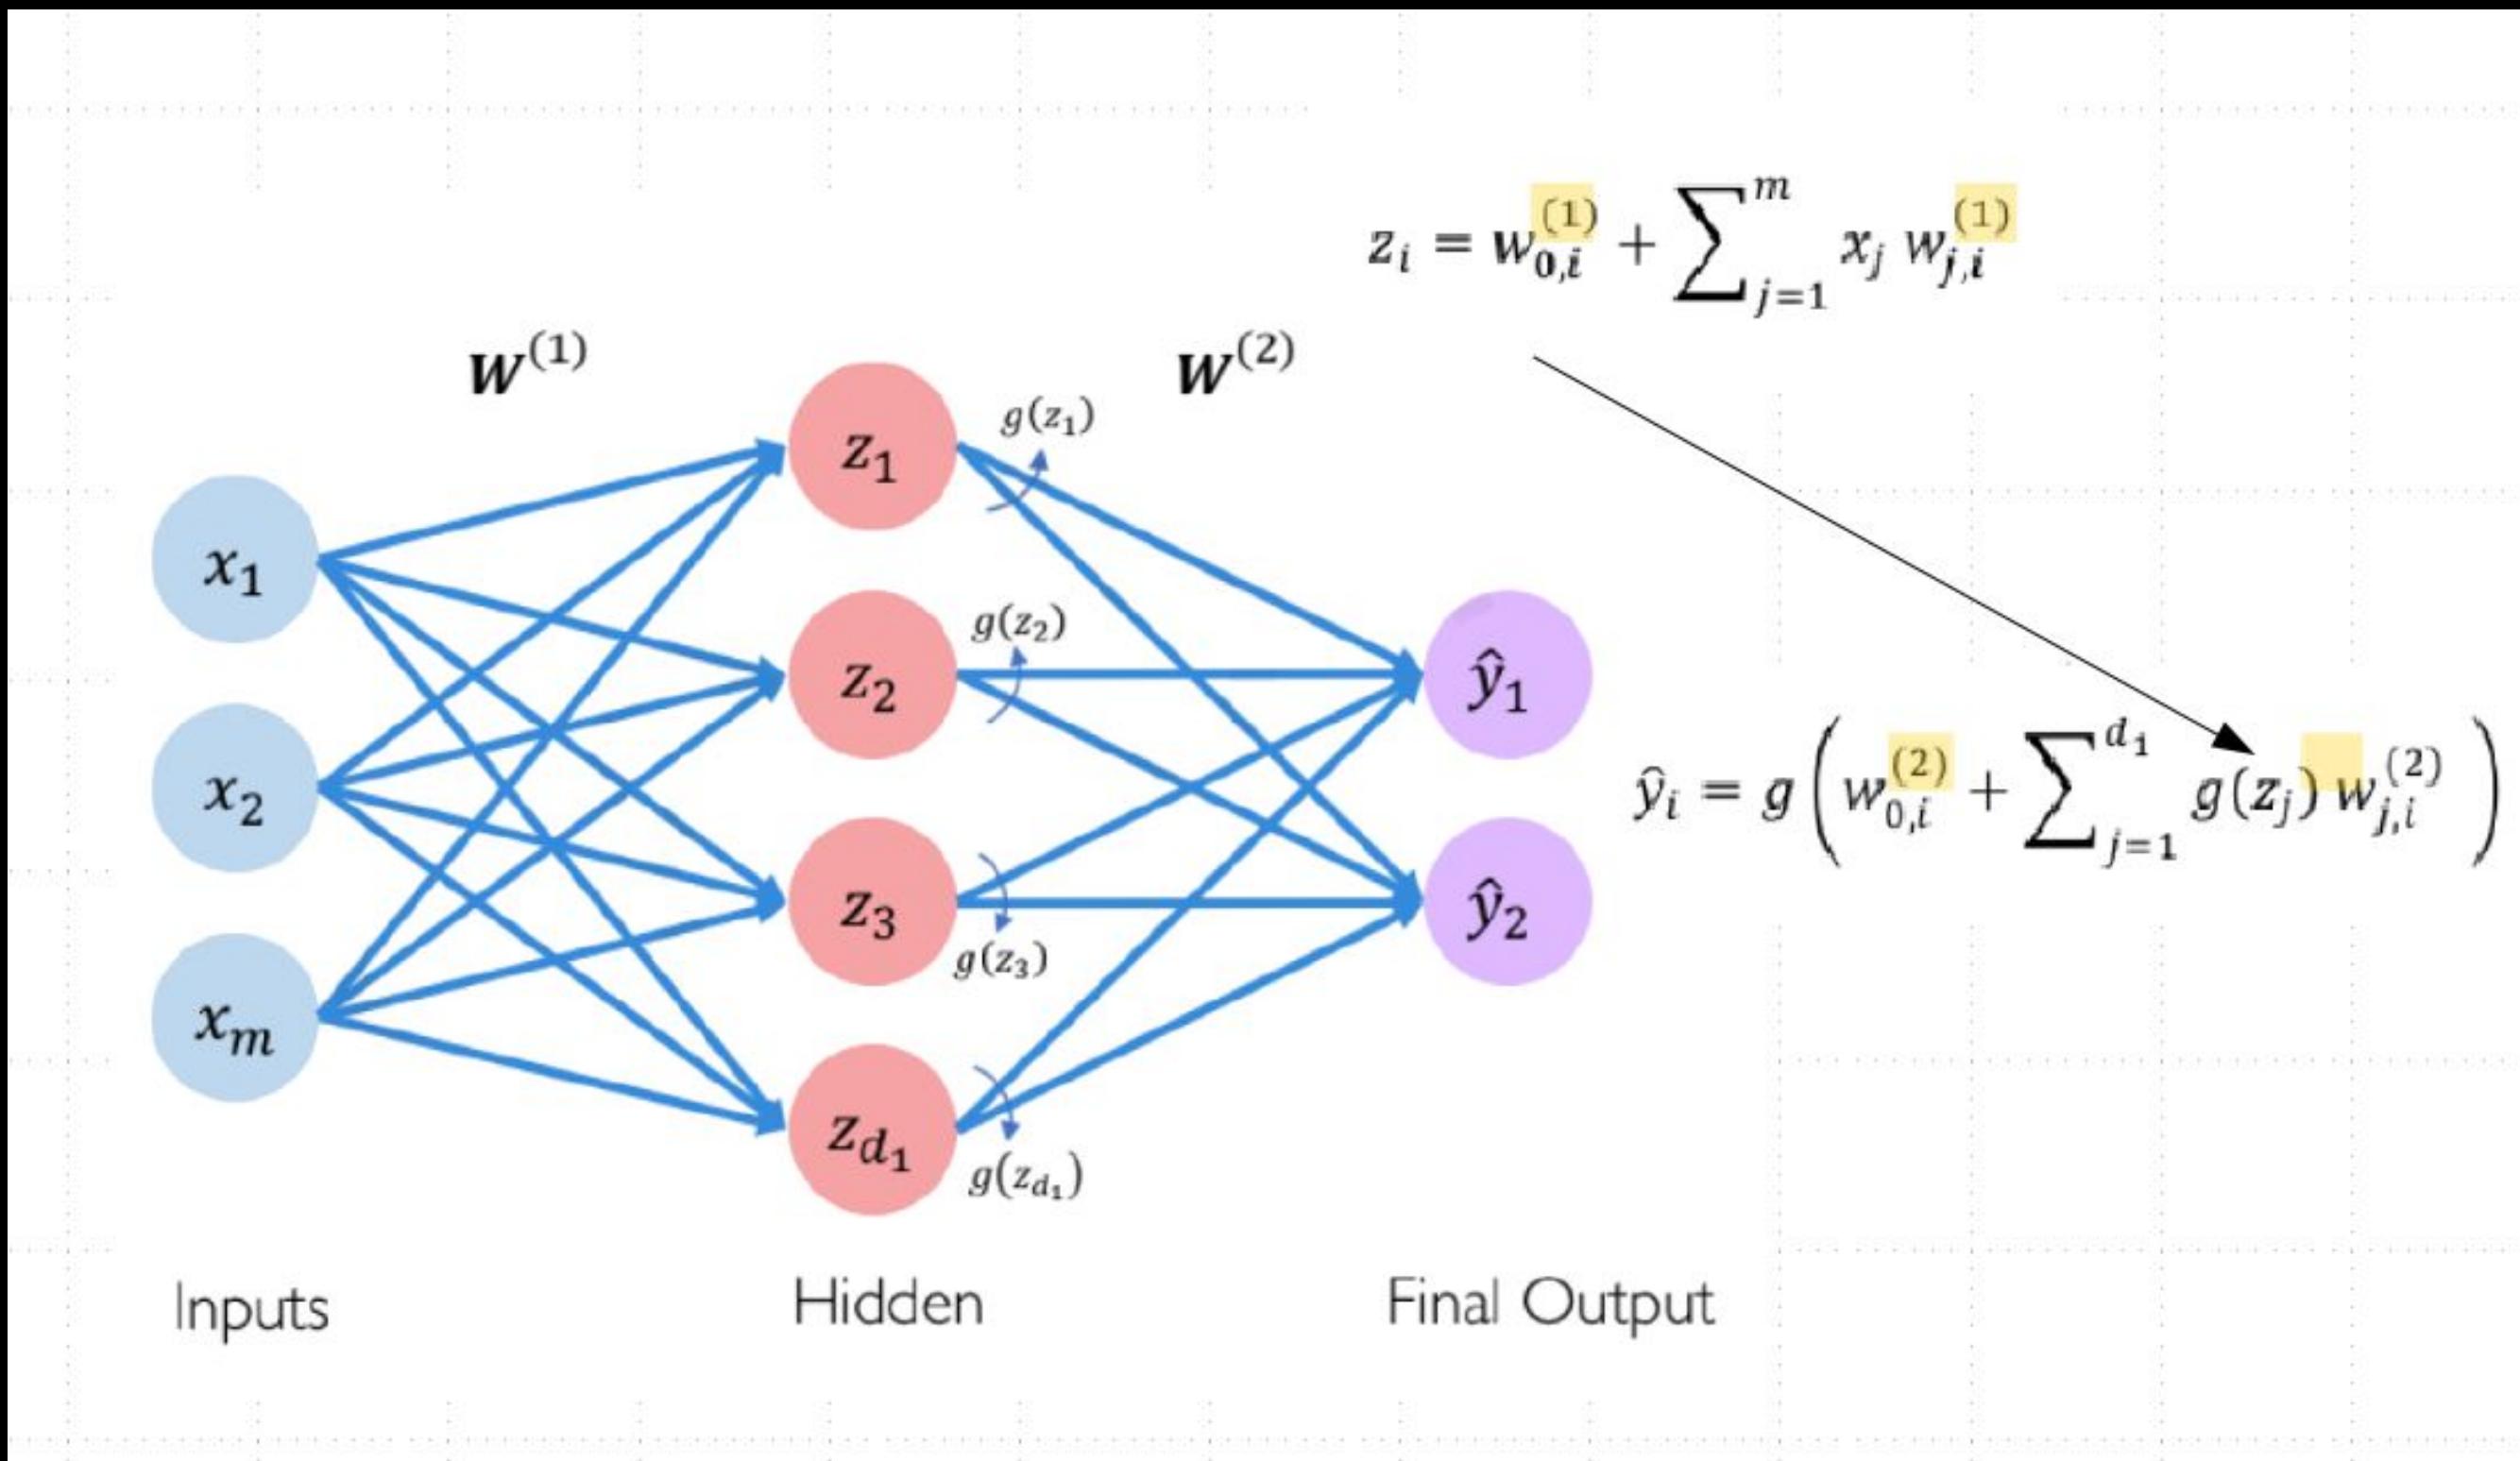
- Neural Network > compuesta por multiples perceptrones (n inputs → m salida)



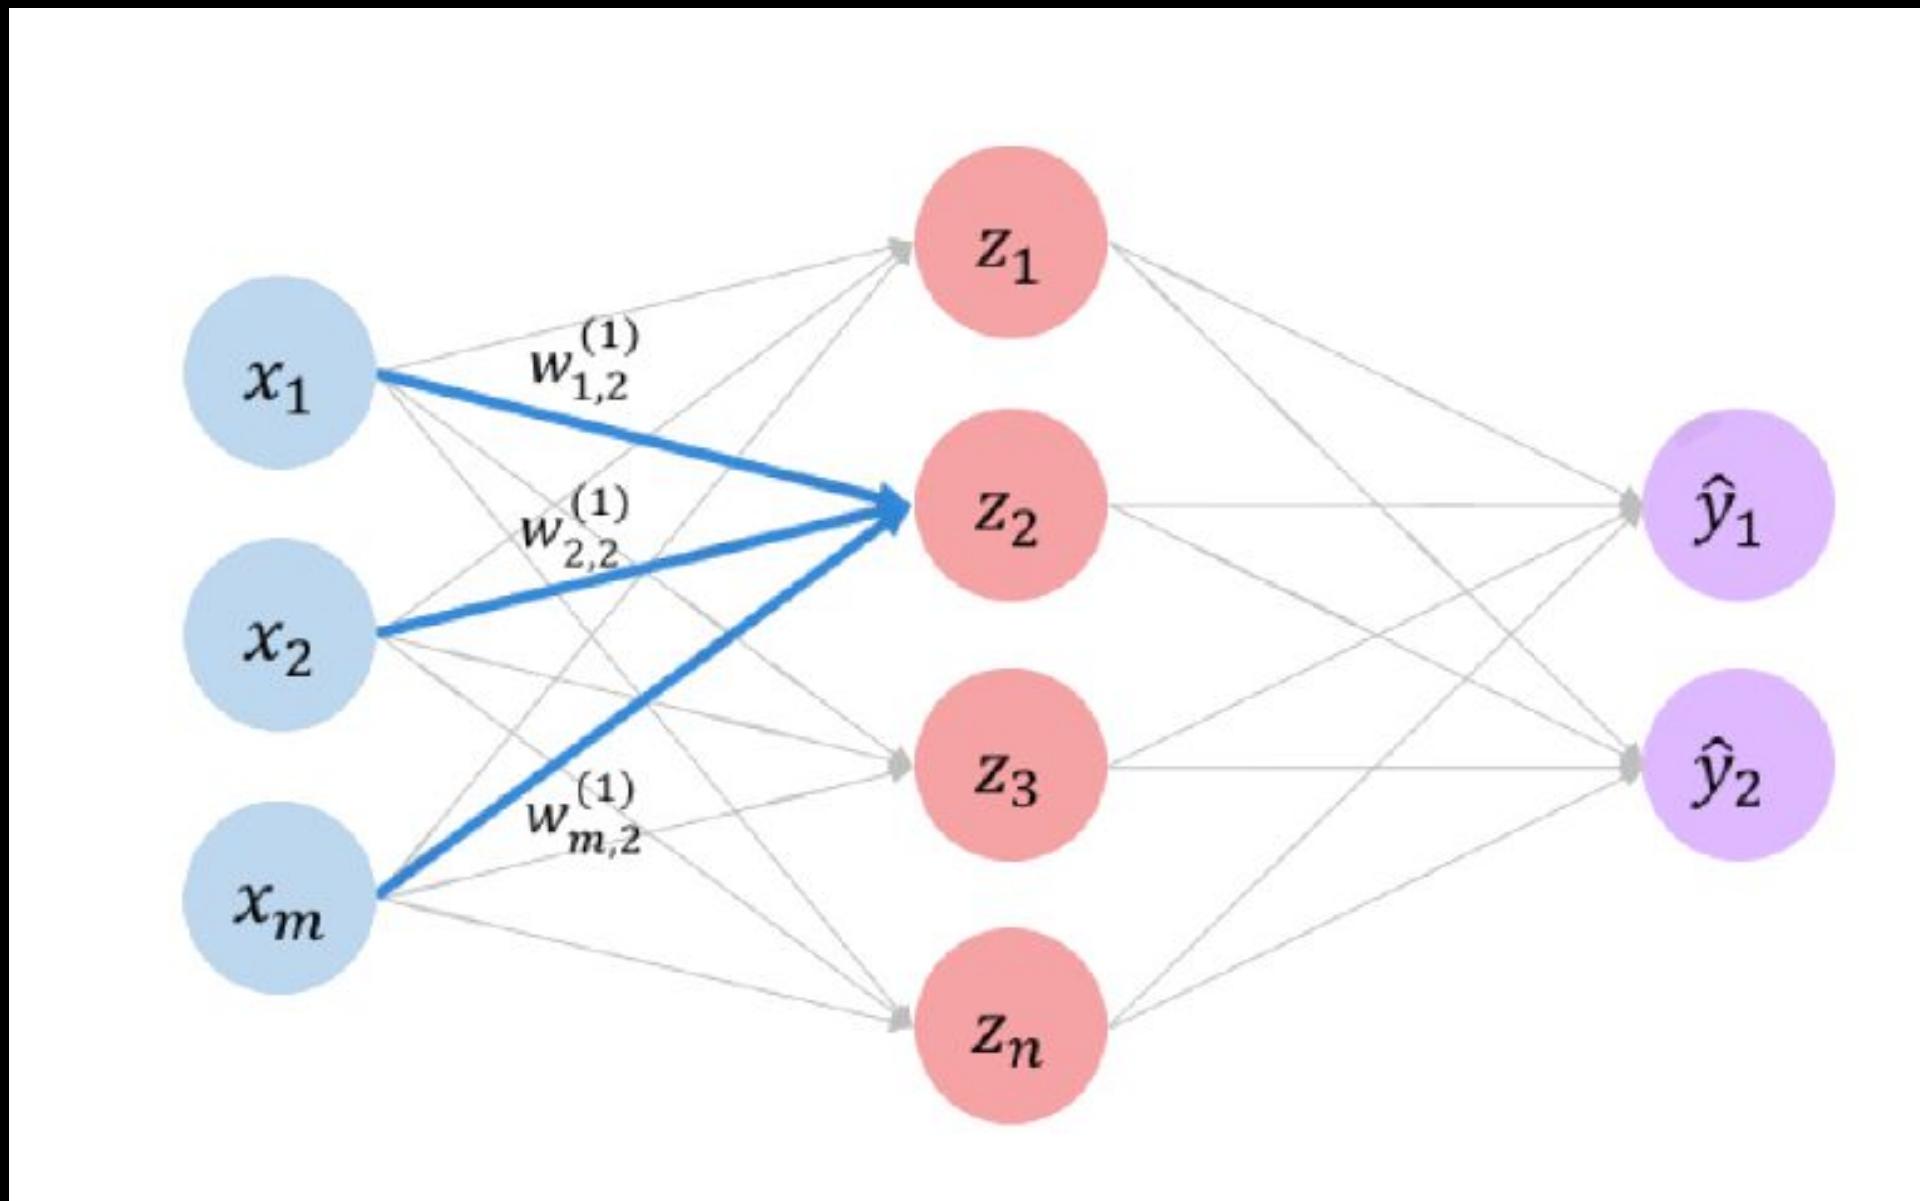
$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

CAPA DENSA: cada entrada está conectada a todas las salidas

1- Red neuronal de una capa



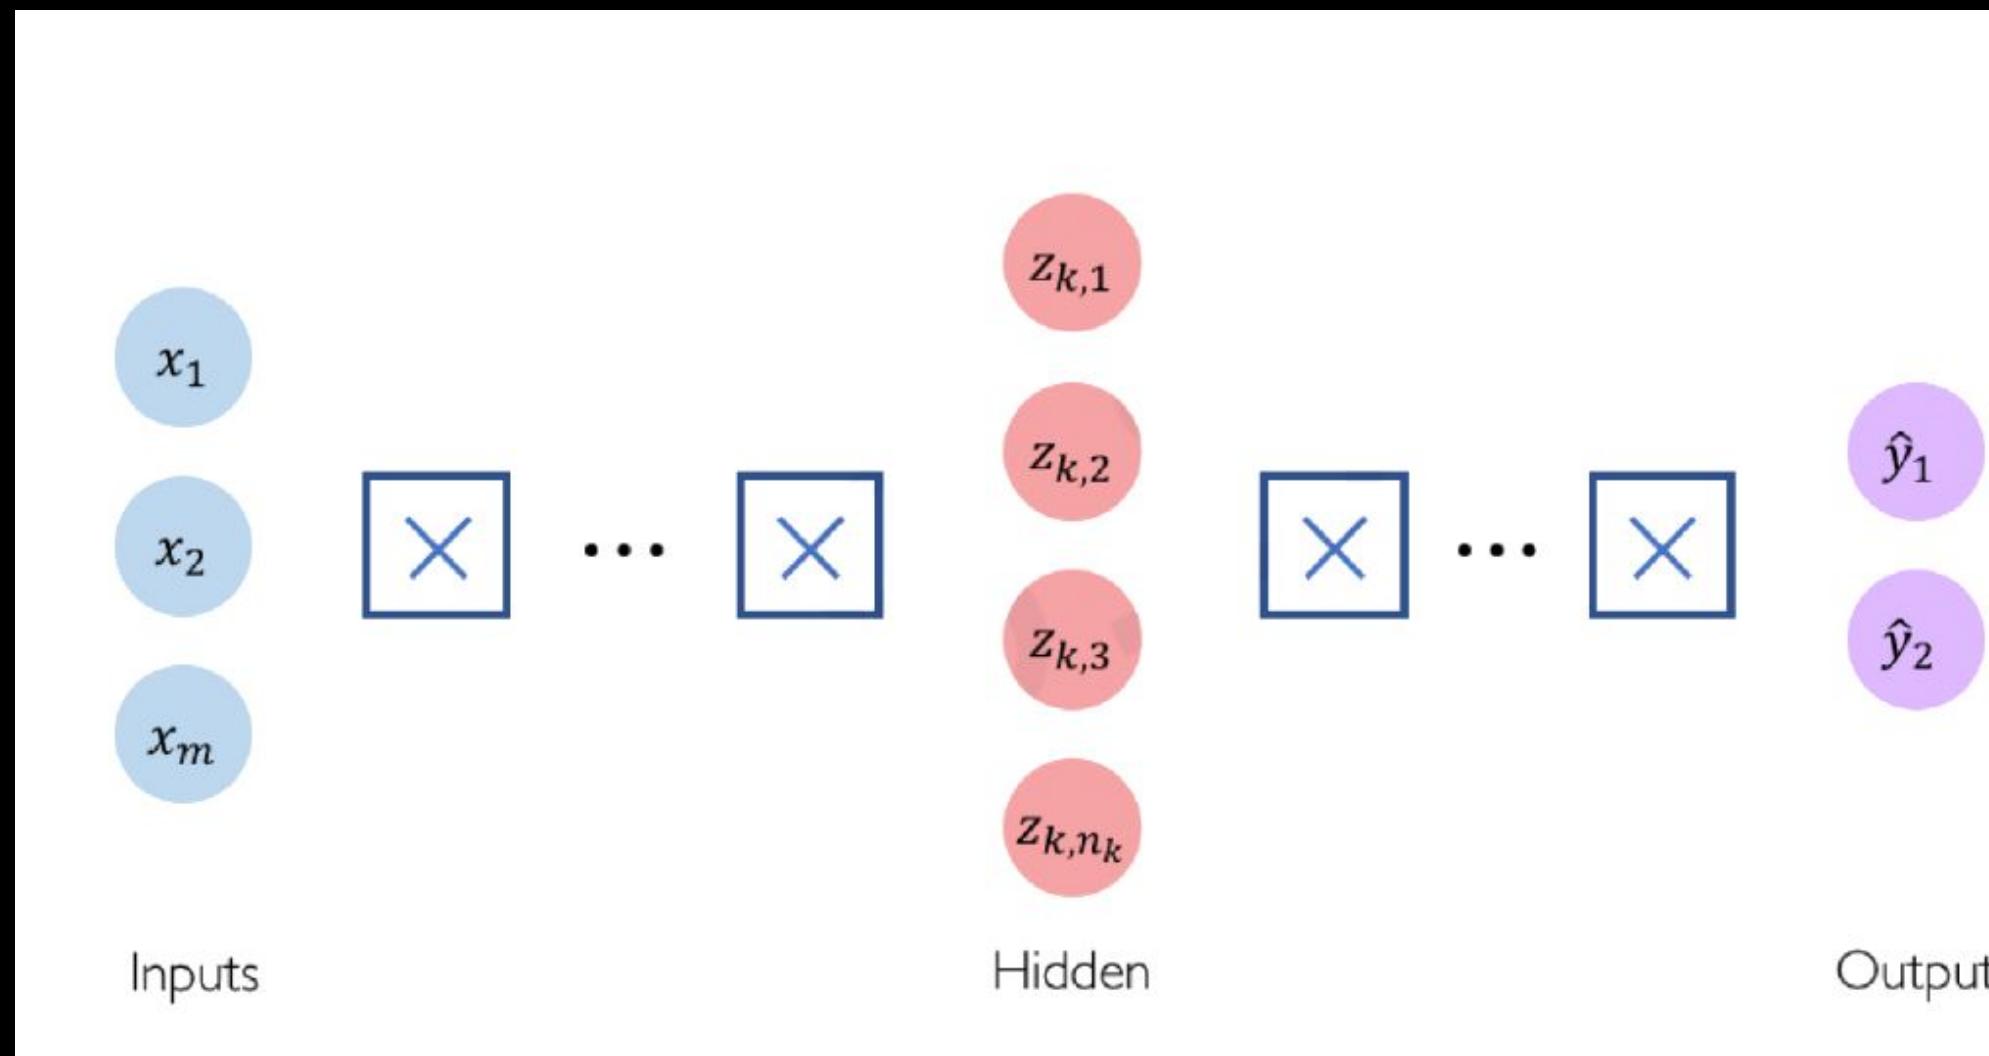
1- Red neuronal de una capa



$$\begin{aligned} z_2 &= w_{0,2}^{(1)} + \sum_{j=1}^m x_j w_{j,2}^{(1)} \\ &= w_{0,2}^{(1)} + x_1 w_{1,2}^{(1)} + x_2 w_{2,2}^{(1)} + x_m w_{m,2}^{(1)} \end{aligned}$$

M-layers Neural Network (red neuronal multi-capa)

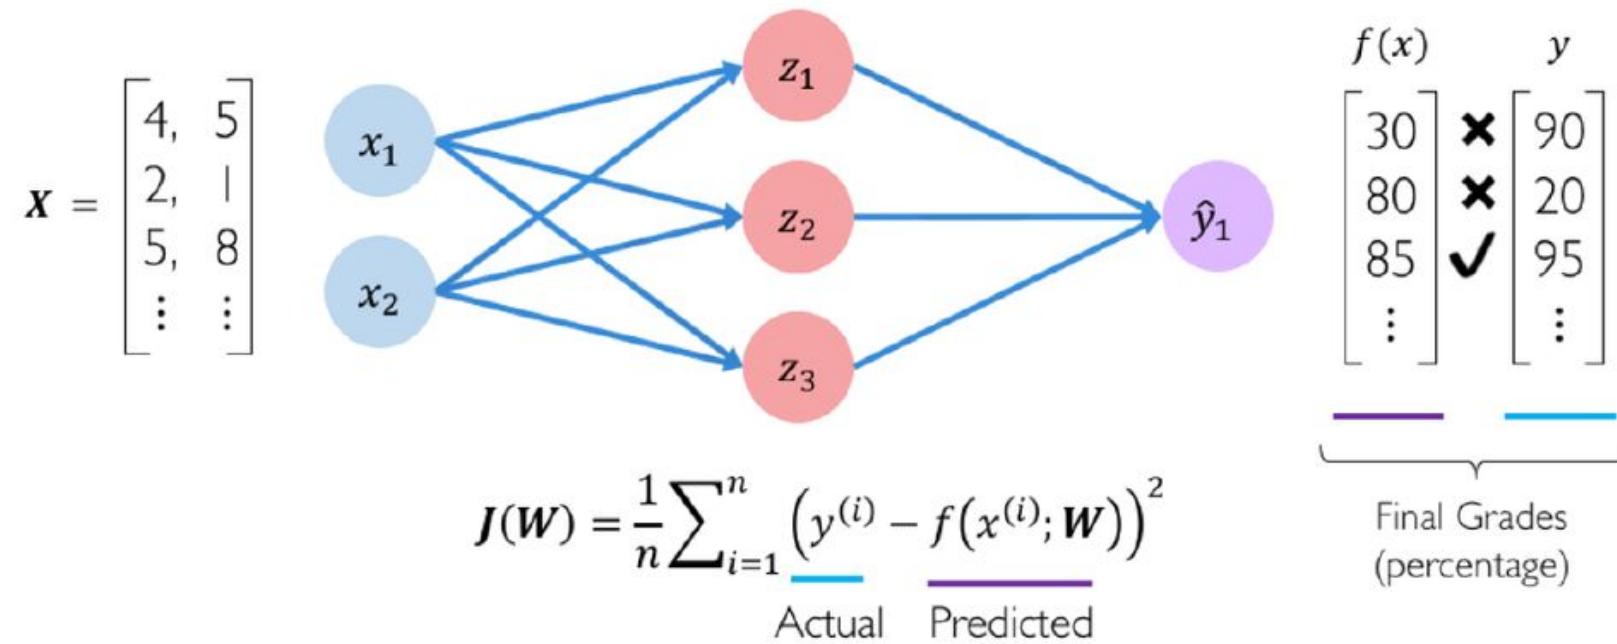
- Deep Neural Network (red neuronal profunda)



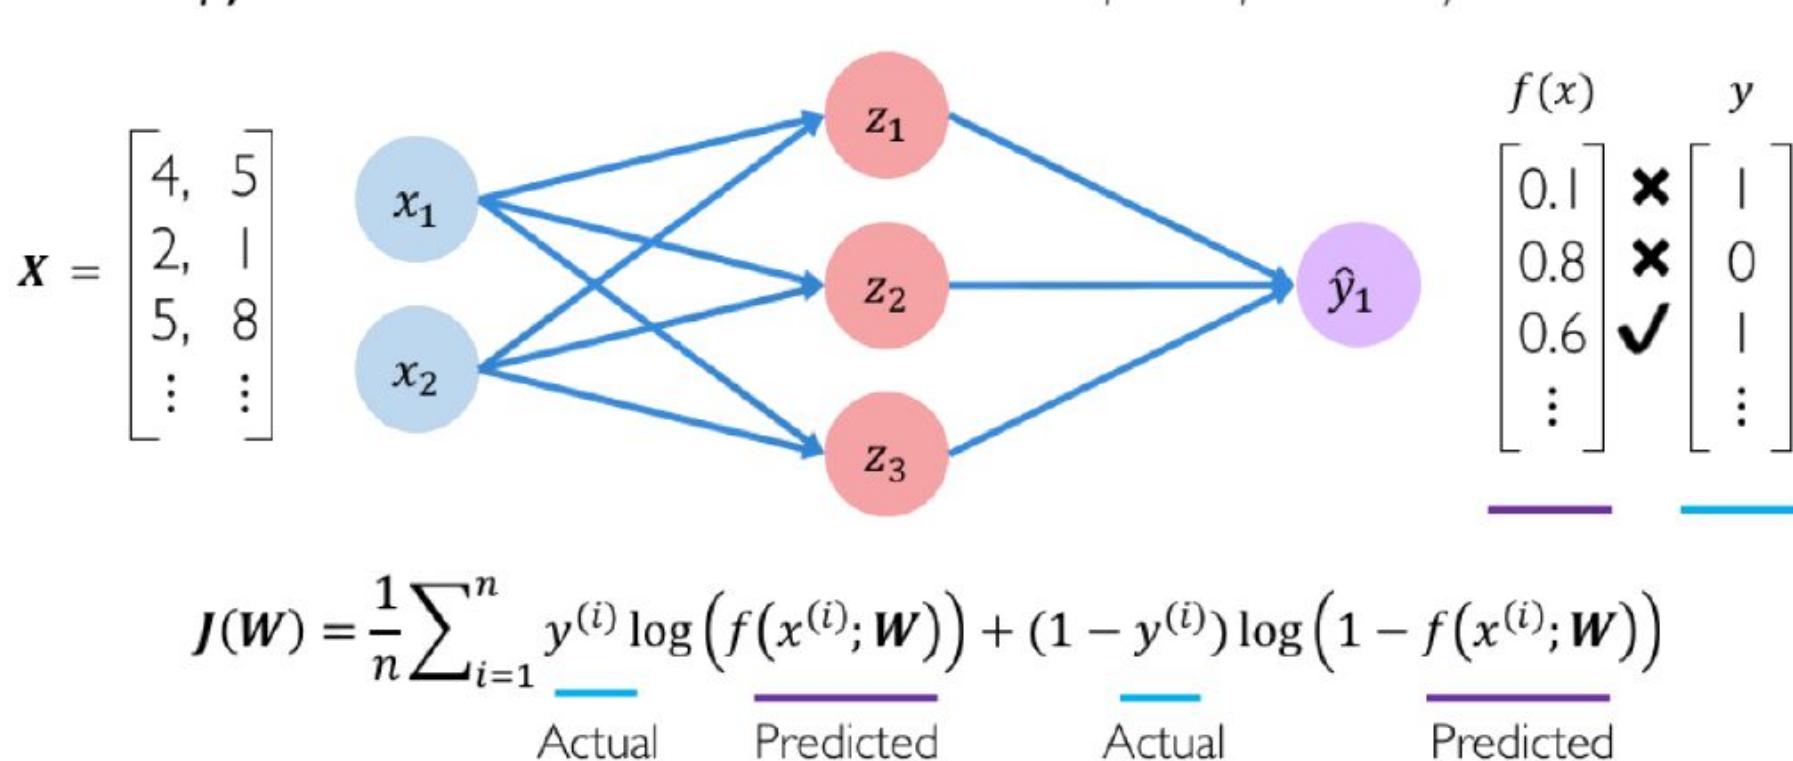
$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{n_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$

Loss Function / Funcion de perdida

Mean squared error loss can be used with regression models that output continuous real numbers



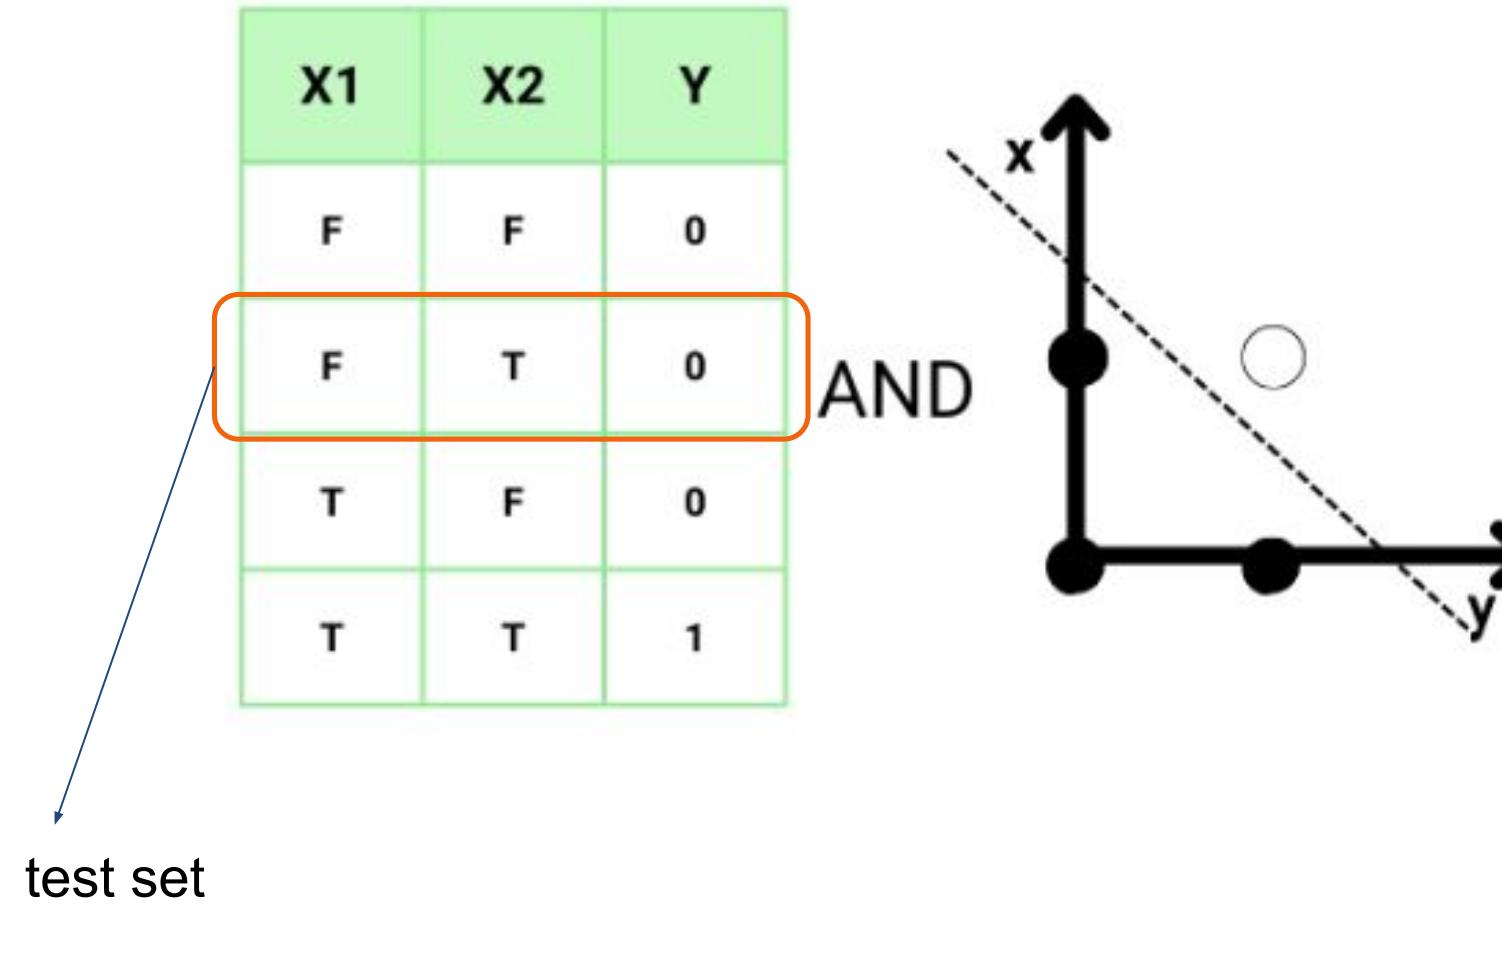
Cross entropy loss can be used with models that output a probability between 0 and 1



symbol	name	equation
\mathcal{L}_1	L_1 loss	$\ y - o\ _1$
\mathcal{L}_2	L_2 loss	$\ y - o\ _2^2$
$\mathcal{L}_1 \circ \sigma$	expectation loss	$\ y - \sigma(o)\ _1$
$\mathcal{L}_2 \circ \sigma$	regularised expectation loss ¹	$\ y - \sigma(o)\ _2^2$
$\mathcal{L}_\infty \circ \sigma$	Chebyshev loss	$\max_j \sigma(o)^{(j)} - y^{(j)} $
hinge	hinge [13] (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{y}^{(j)} o^{(j)})$
hinge ²	squared hinge (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{y}^{(j)} o^{(j)})^2$
hinge ³	cubed hinge (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{y}^{(j)} o^{(j)})^3$
log	log (cross entropy) loss	$-\sum_j y^{(j)} \log \sigma(o)^{(j)}$
log ²	squared log loss	$-\sum_j [y^{(j)} \log \sigma(o)^{(j)}]^2$
tan	Tanimoto loss	$-\sum_j \frac{\sigma(o)^{(j)} y^{(j)}}{\ \sigma(o)\ _2^2 + \ y\ _2^2 - \sum_j \sigma(o)^{(j)} y^{(j)}}$
D _{CS}	Cauchy-Schwarz Divergence [3]	$-\log \frac{\sum_j \sigma(o)^{(j)} y^{(j)}}{\ \sigma(o)\ _2 \ y\ _2}$

- <https://arxiv.org/pdf/1702.05659.pdf>

Entrenemos un perceptrón



Recordemos

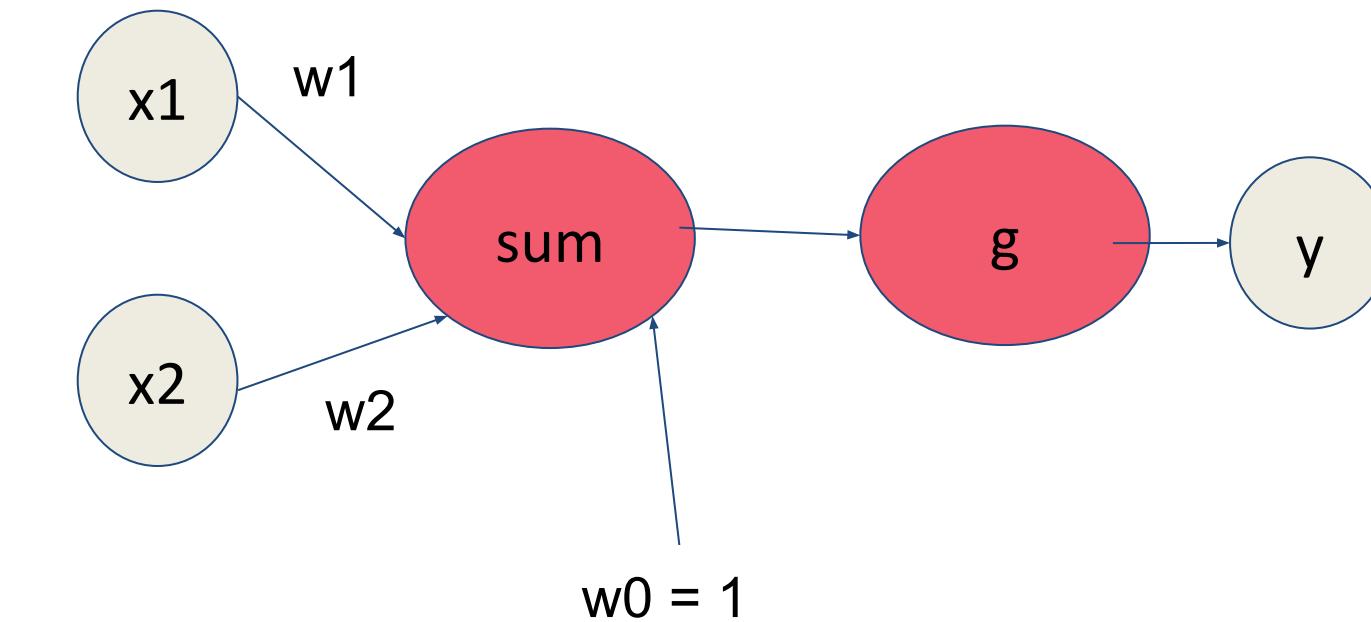
Linear combination of inputs

$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function

Bias

Nuestra arquitectura sería :

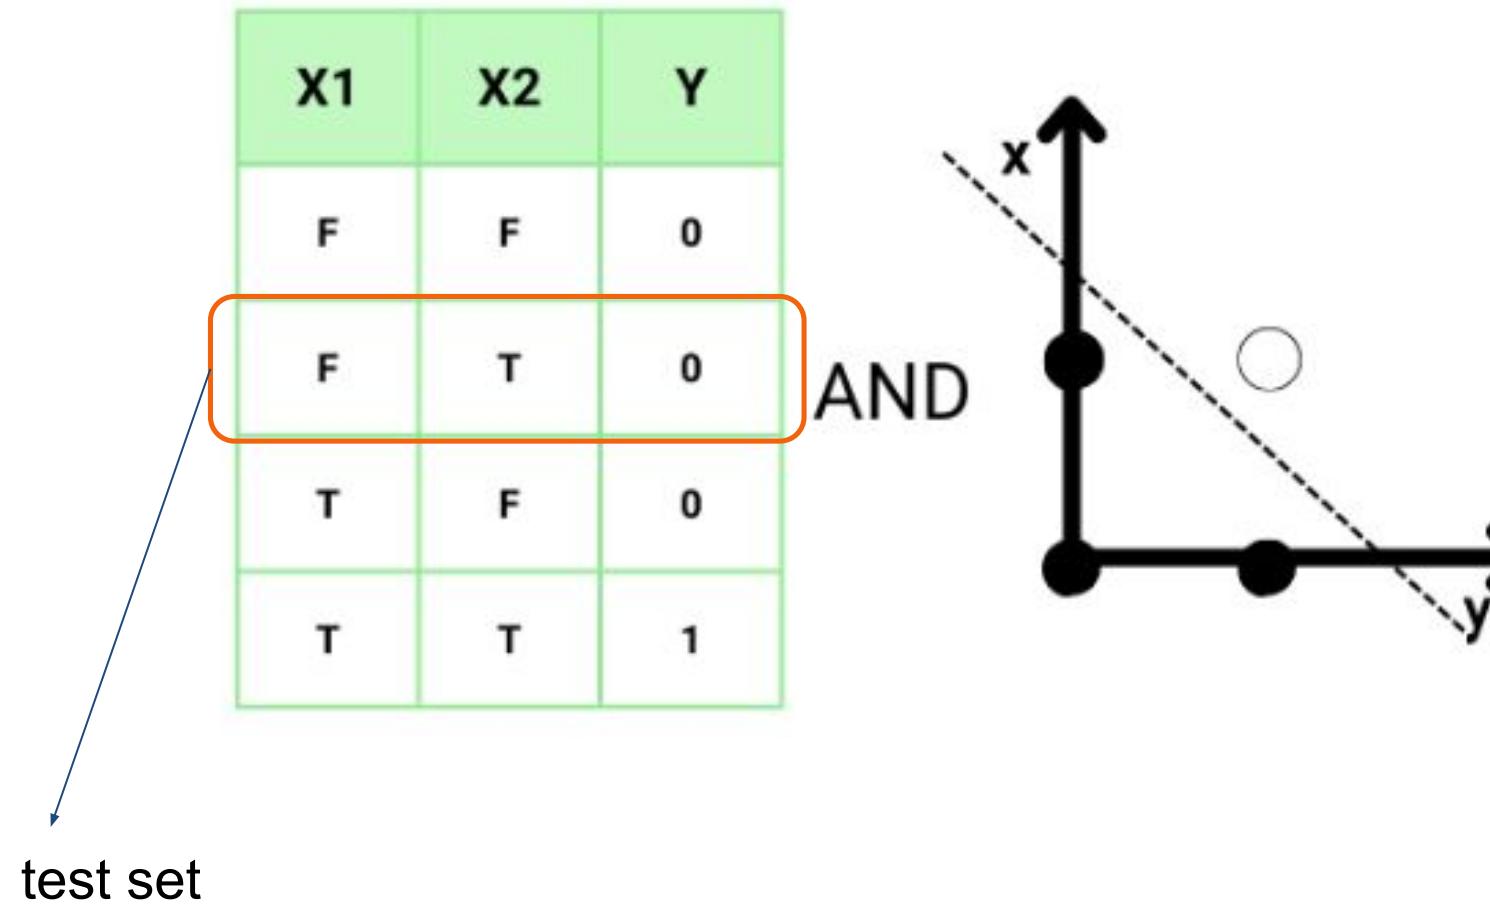


Cuanto vale g ?

Supongamos que g devuelve 1 si es > 1.5 y devuelve 0 en caso contrario

Entrenemos un perceptrón

Comienza el entrenamiento It =1



se obtienen de manera aleatoria los valores de los pesos.

Supongamos $w_1=0.9$ y $w_2=0.9$

Entonces.

$$y_1 = g(0.9 * 0 + 0.9 * 0 + 1) = g(1) = 1 \text{ pero } y(0,0) = 0$$

$$e = \text{abs}(0-0)=0$$

$$y_2 = g(0.9 * 1 + 0.9 * 0 + 1) = g(1.9) = 1 \text{ pero } y(1,0) = 0$$

$$e = \text{abs}(1.9-0)=1.9$$

$$y_3 = g(0.9 * 1 + 0.9 * 1 + 1) = g(2.8) = 1 \text{ y ademas } y(1,1) = 1$$

$$e = 0$$

Hasta aquí el modelo generado tiene la siguiente forma:

$$y_{\text{aprox}} = g(0.9 x_1 + 0.9 x_2 + 1)$$

Como seria la predicción del test set?

$$y = g(0.9 * 1 + 0.9 * 0 + 1) = g(1.9) = 1 \text{ pero deberia ser } 0!!!!$$

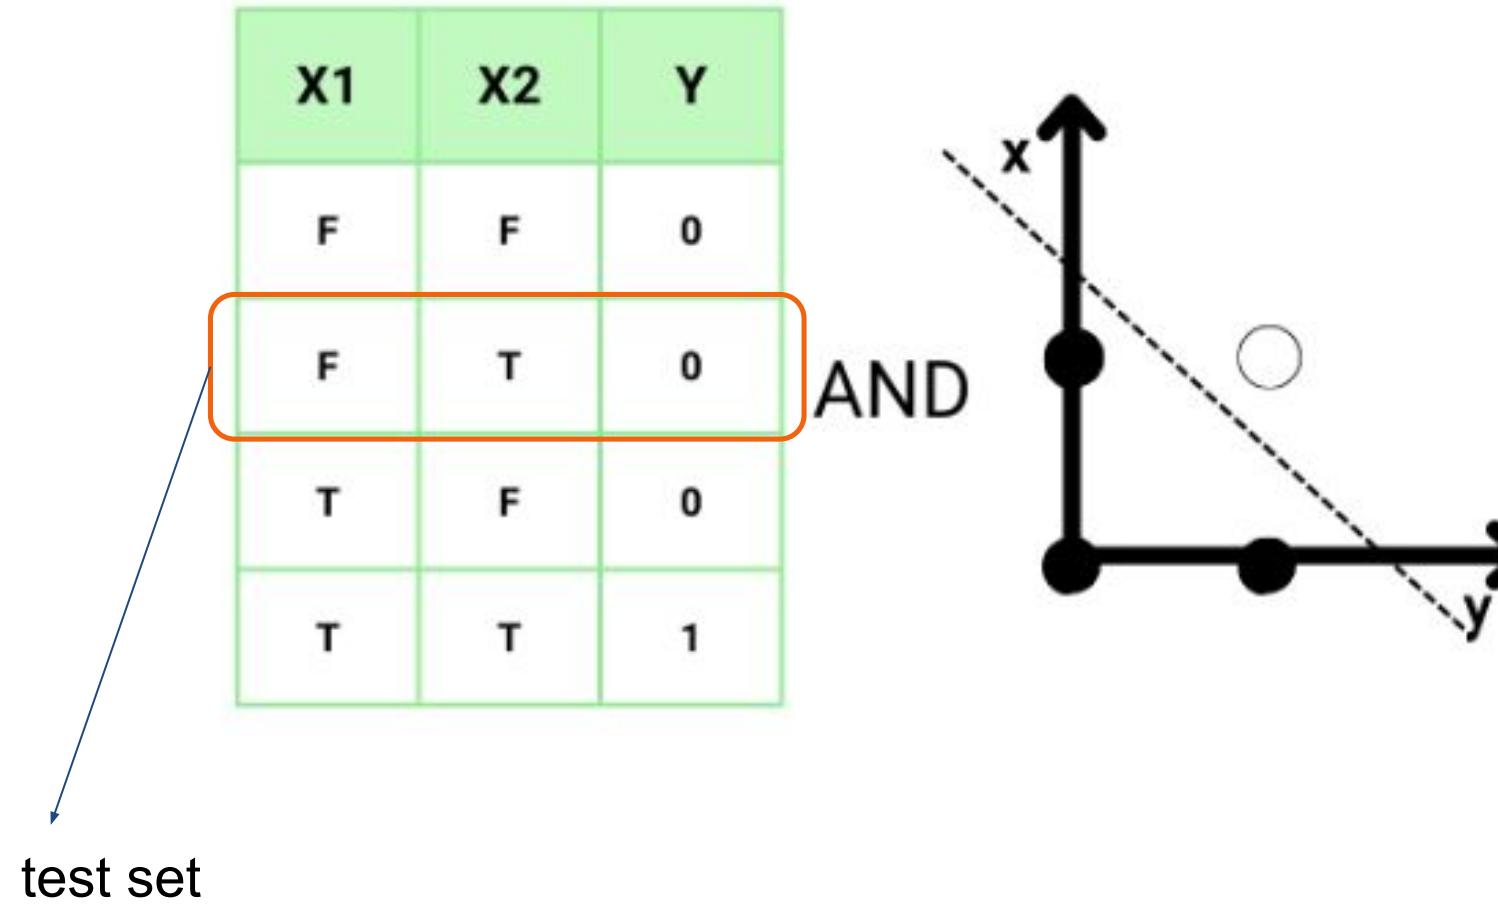
$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Annotations:

- Output (purple arrow pointing down)
- Non-linear activation function (yellow arrow pointing up)
- Linear combination of inputs (red arrow pointing down)
- Bias (green arrow pointing up)

Entrenemos un perceptrón

Seguimos entrenando, $l = 2$



se obtienen de manera aleatoria los valores de los pesos.

Supongamos $w_1=0.4$ y $w_2=0.4$

Entonces:

$$y_1 = g(0.4 \cdot 0 + 0.4 \cdot 0 + 1) = g(1) = 1 \text{ pero } y(0,0) = 0$$

$$e = \text{abs}(0-0) = 0$$

$$y_2 = g(0.4 \cdot 1 + 0.4 \cdot 0 + 1) = g(1.4) = 0 \text{ pero } y(1,0) = 0$$

$$e = \text{abs}(0-0) = 0$$

$$y_3 = g(0.9 \cdot 1 + 0.9 \cdot 1 + 1) = g(2.8) = 1 \text{ y ademas } y(1,1) = 1$$

$$e = 0$$

Hasta aquí el modelo generado tiene la siguiente forma:

$$y_{\text{aprox}} = g(0.4 \cdot x_1 + 0.4 \cdot x_2 + 1)$$

Como sería la predicción del test set?

$$y = g(0.4 \cdot 0 + 0.4 \cdot 1 + 1) = g(1.4) = 0 \text{ y el resultado debía ser } 1!!! :)$$

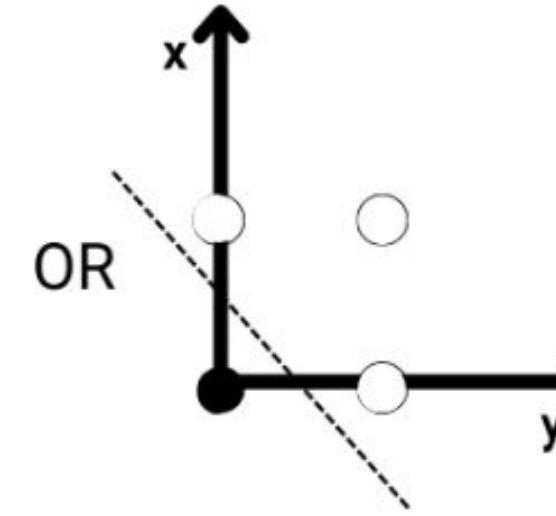
$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Annotations:

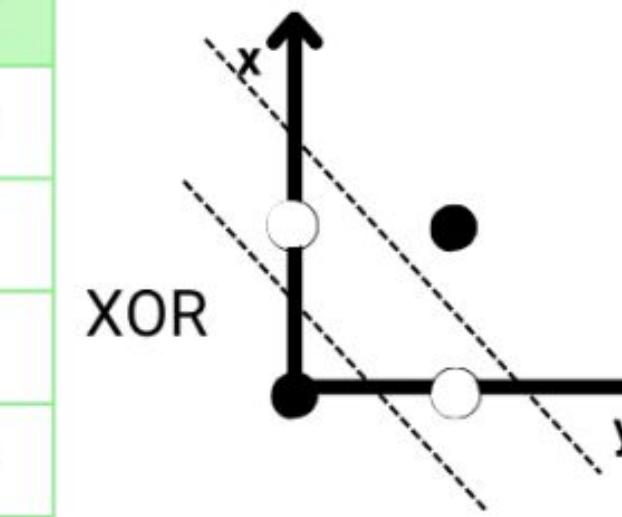
- Output (purple arrow)
- Linear combination of inputs (red arrow)
- Non-linear activation function (yellow arrow)
- Bias (green arrow)

Entrenemos un perceptrón

X1	X2	Y
F	F	0
F	T	1
T	F	1
T	T	1



X1	X2	Y
F	F	0
F	T	1
T	F	1
T	T	0



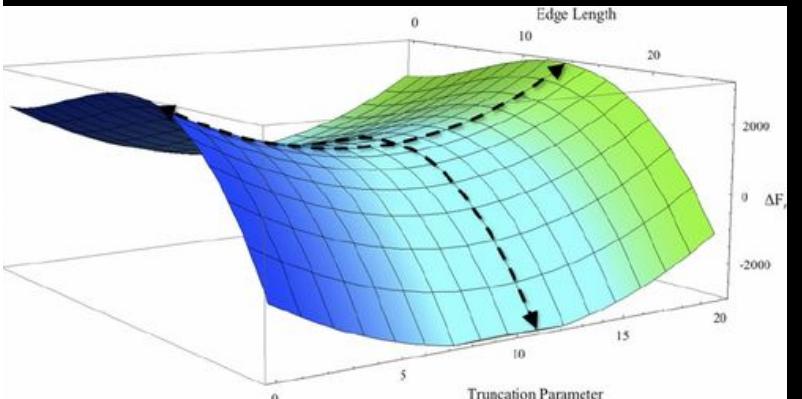
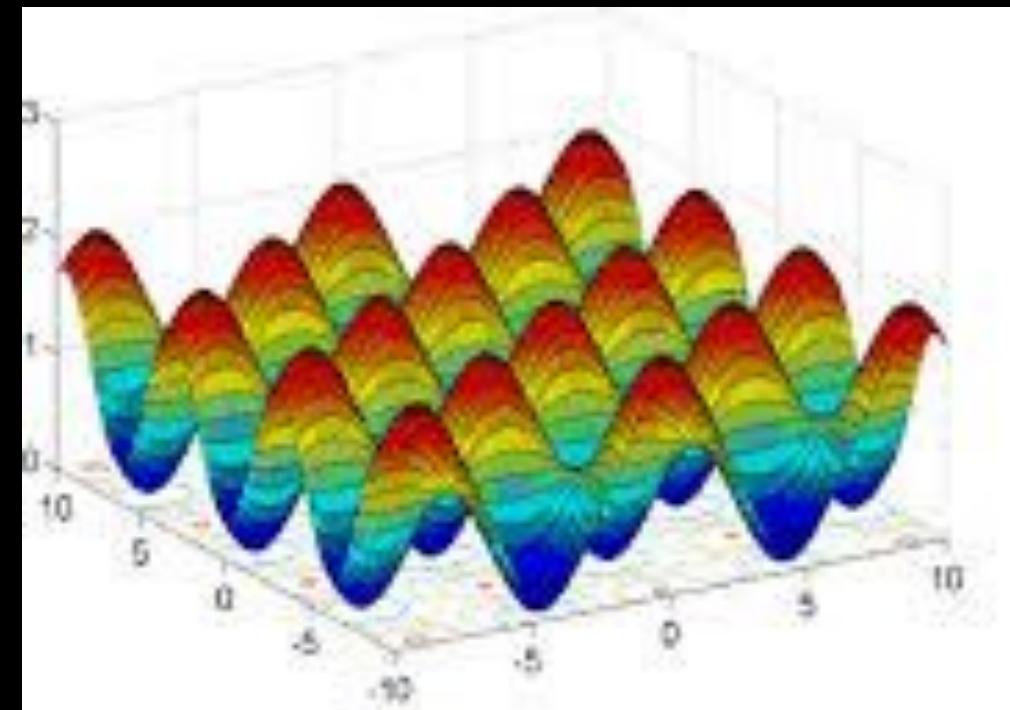
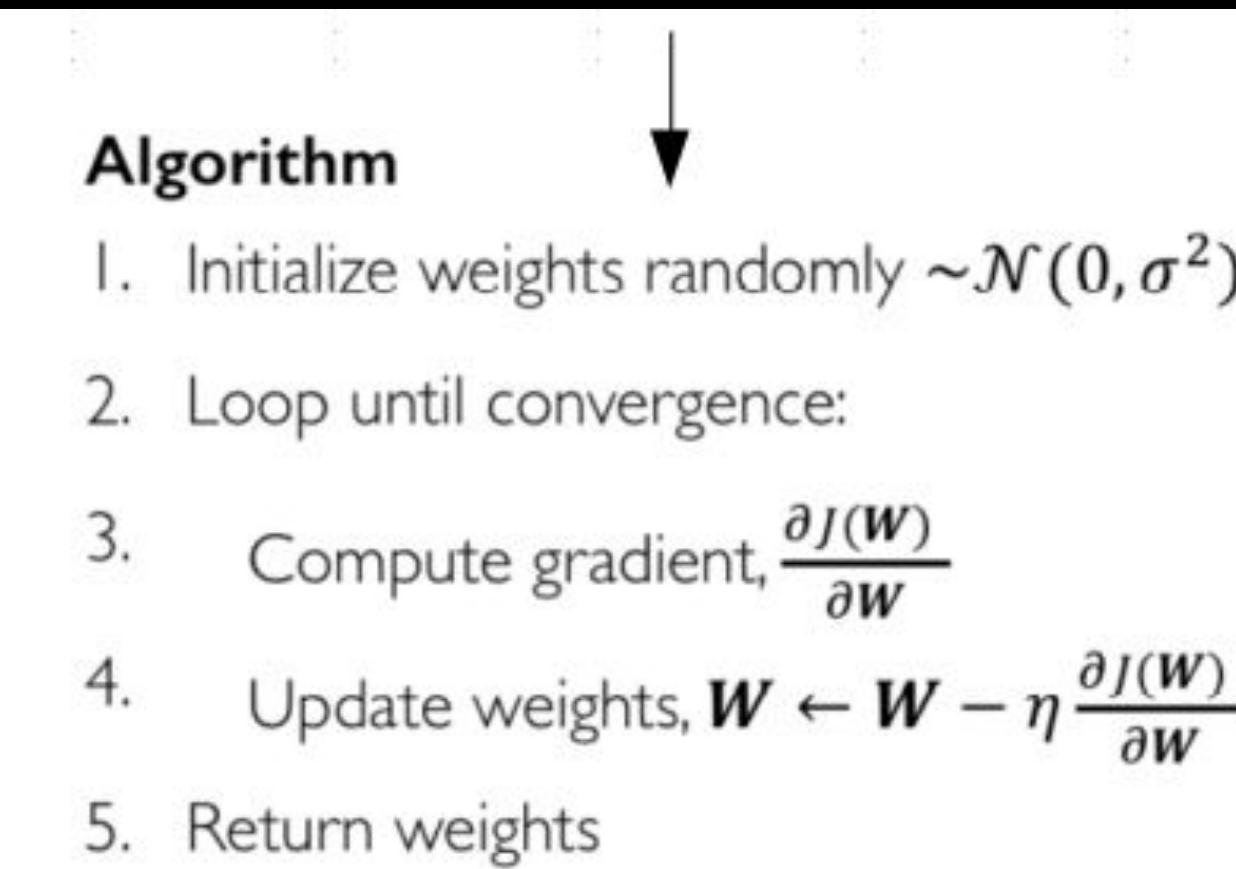
pensar estos dos casos!

Training (optimización de la función de pérdida)

- encontrar los valores de los pesos para que la función de pérdida (loss) o costo sea la mínima posible
- Problema de optimización

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$
$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$

METODO DEL GRADIENTE DESCENDIENTE

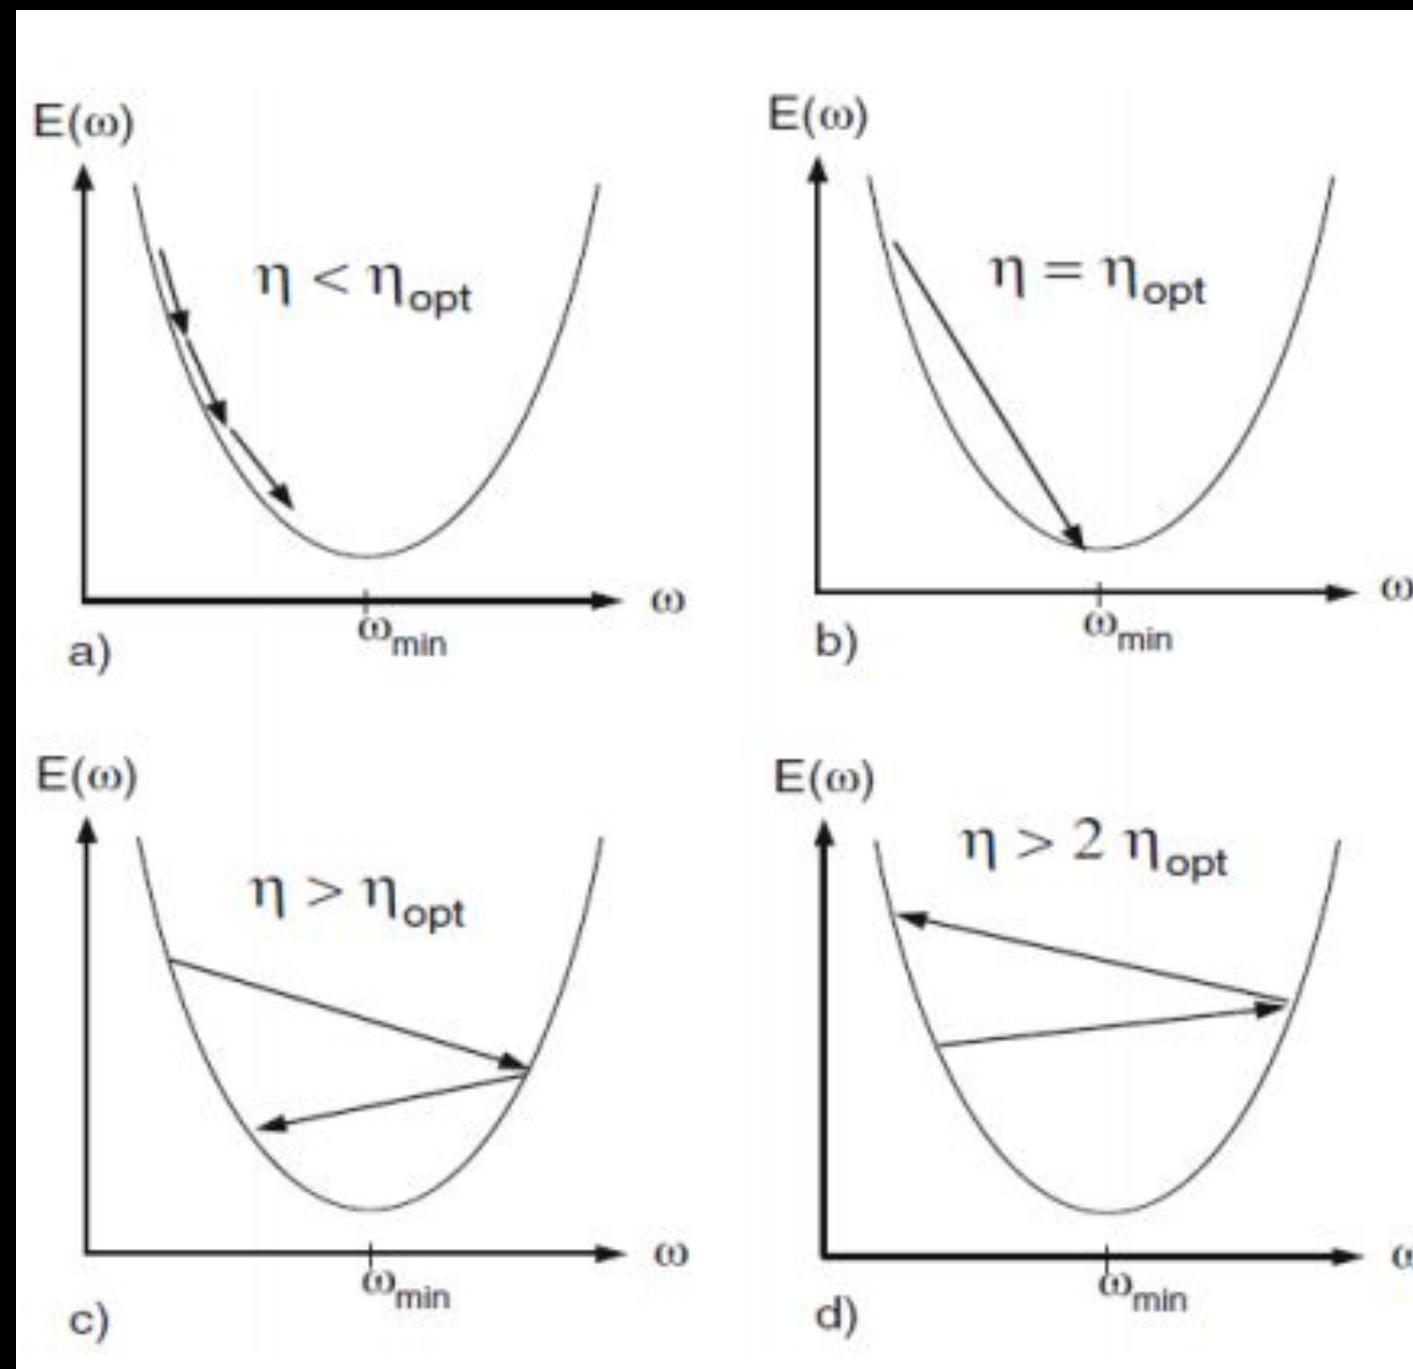


Tiene algunos problemas:

- Minimos locales minima
- Silla de montar: cuando llega a una región de gradiente cero, no puede escapar



Training (optimización de la función de pérdida)



$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$

How can we set the learning rate?

- <https://ruder.io/optimizing-gradient-descent/index.html>

Gradient Descent Algorithms

Algorithm

- SGD
- Adam
- Adadelta
- Adagrad
- RMSProp

TF Implementation

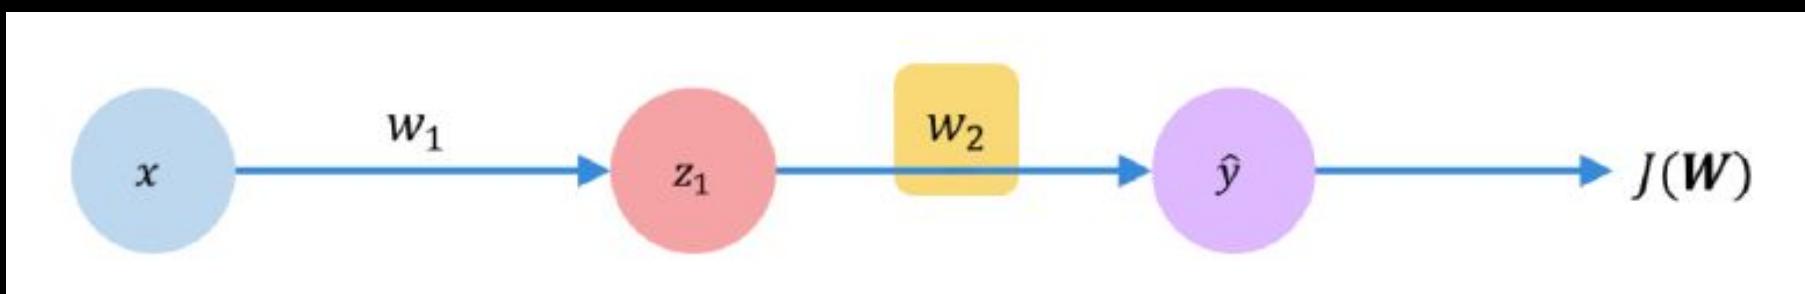
- `tf.keras.optimizers.SGD`
- `tf.keras.optimizers.Adam`
- `tf.keras.optimizers.Adadelta`
- `tf.keras.optimizers.Adagrad`
- `tf.keras.optimizers.RMSProp`

Reference

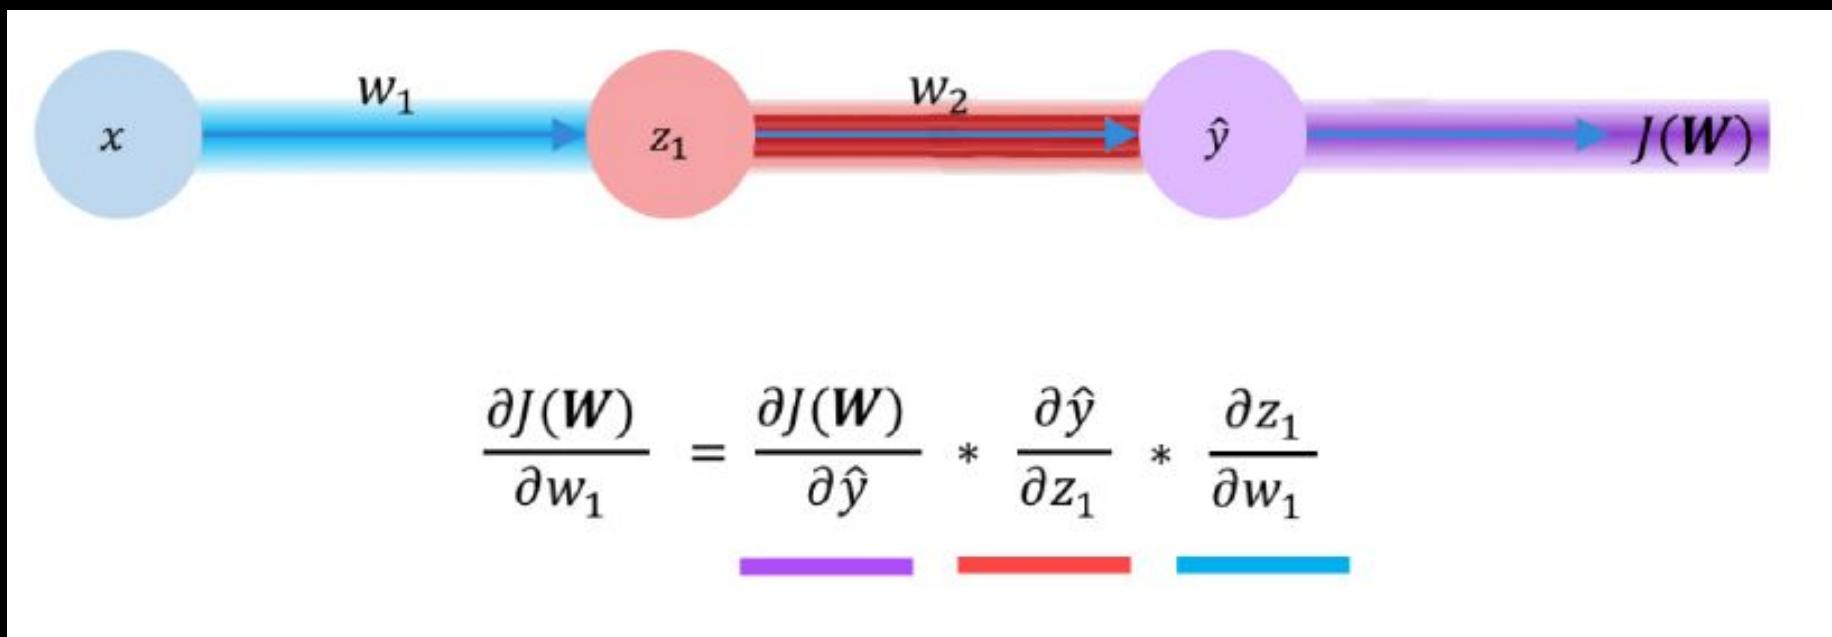
- Kiefer & Wolfowitz. "Stochastic Estimation of the Maximum of a Regression Function." 1952.
- Kingma et al. "Adam: A Method for Stochastic Optimization." 2014.
- Zeiler et al. "ADADELTA: An Adaptive Learning Rate Method." 2012.
- Duchi et al. "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization." 2011.

Backpropagation

- Que tanto afectan los cambios pequeños en uno o más pesos?



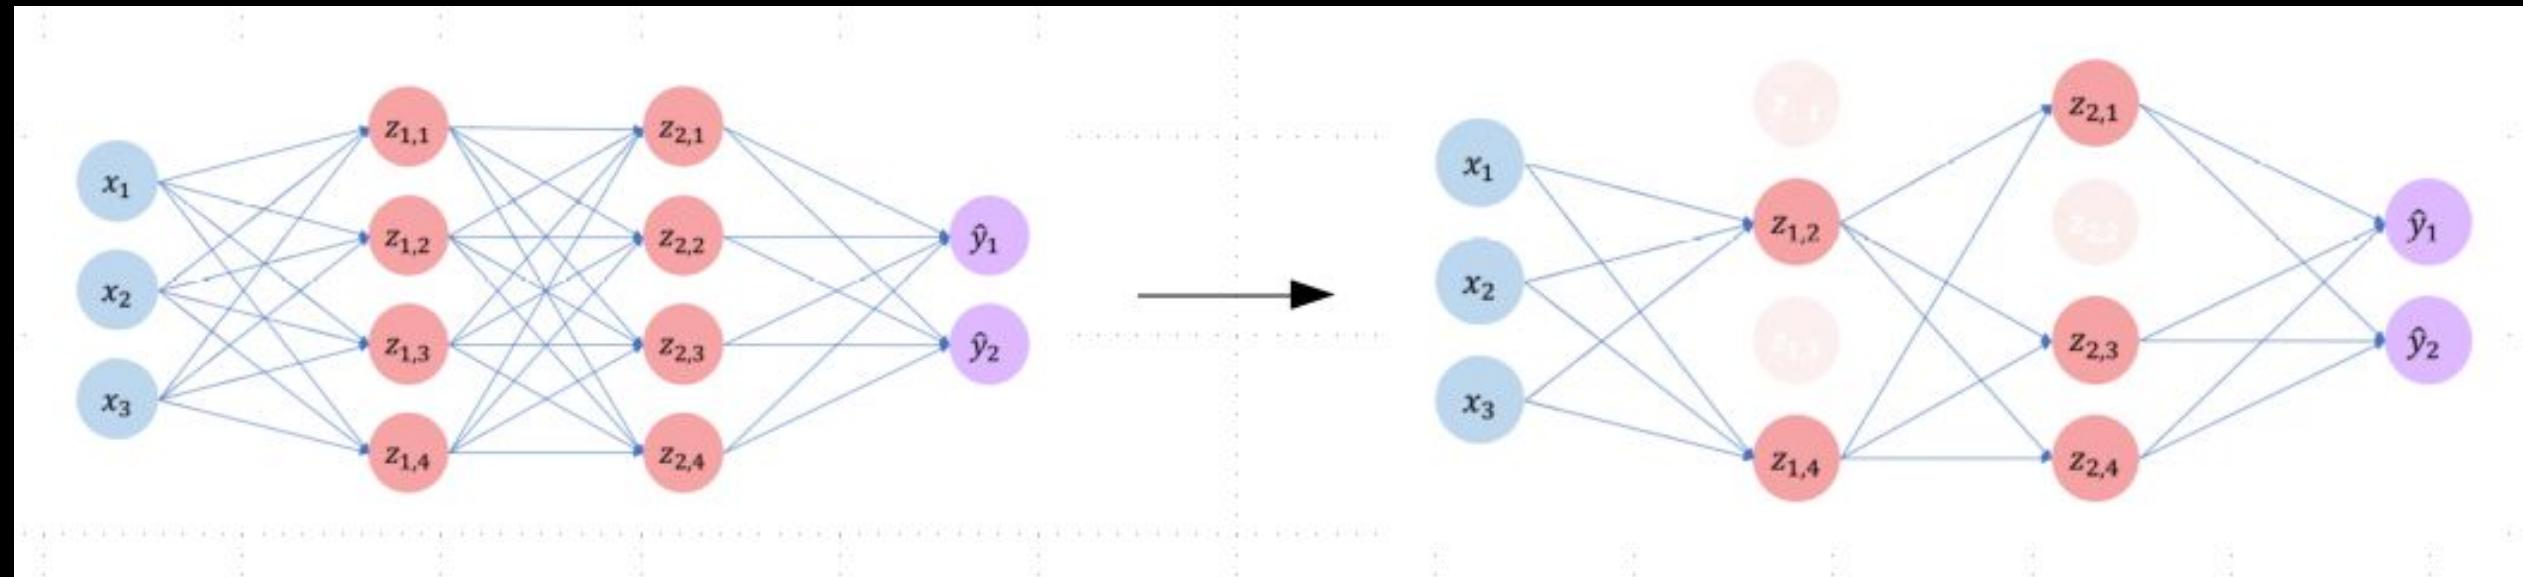
- Se aplica la regla de la cadena



- Repite para cada uno de los pesos de la red usando la forma del gradiente de la siguiente capa!

Regularización

- Técnica que agrega restricciones al problema de optimización (evita modelos complejos)
- Busca lograr modelos más generales



DROP OUT

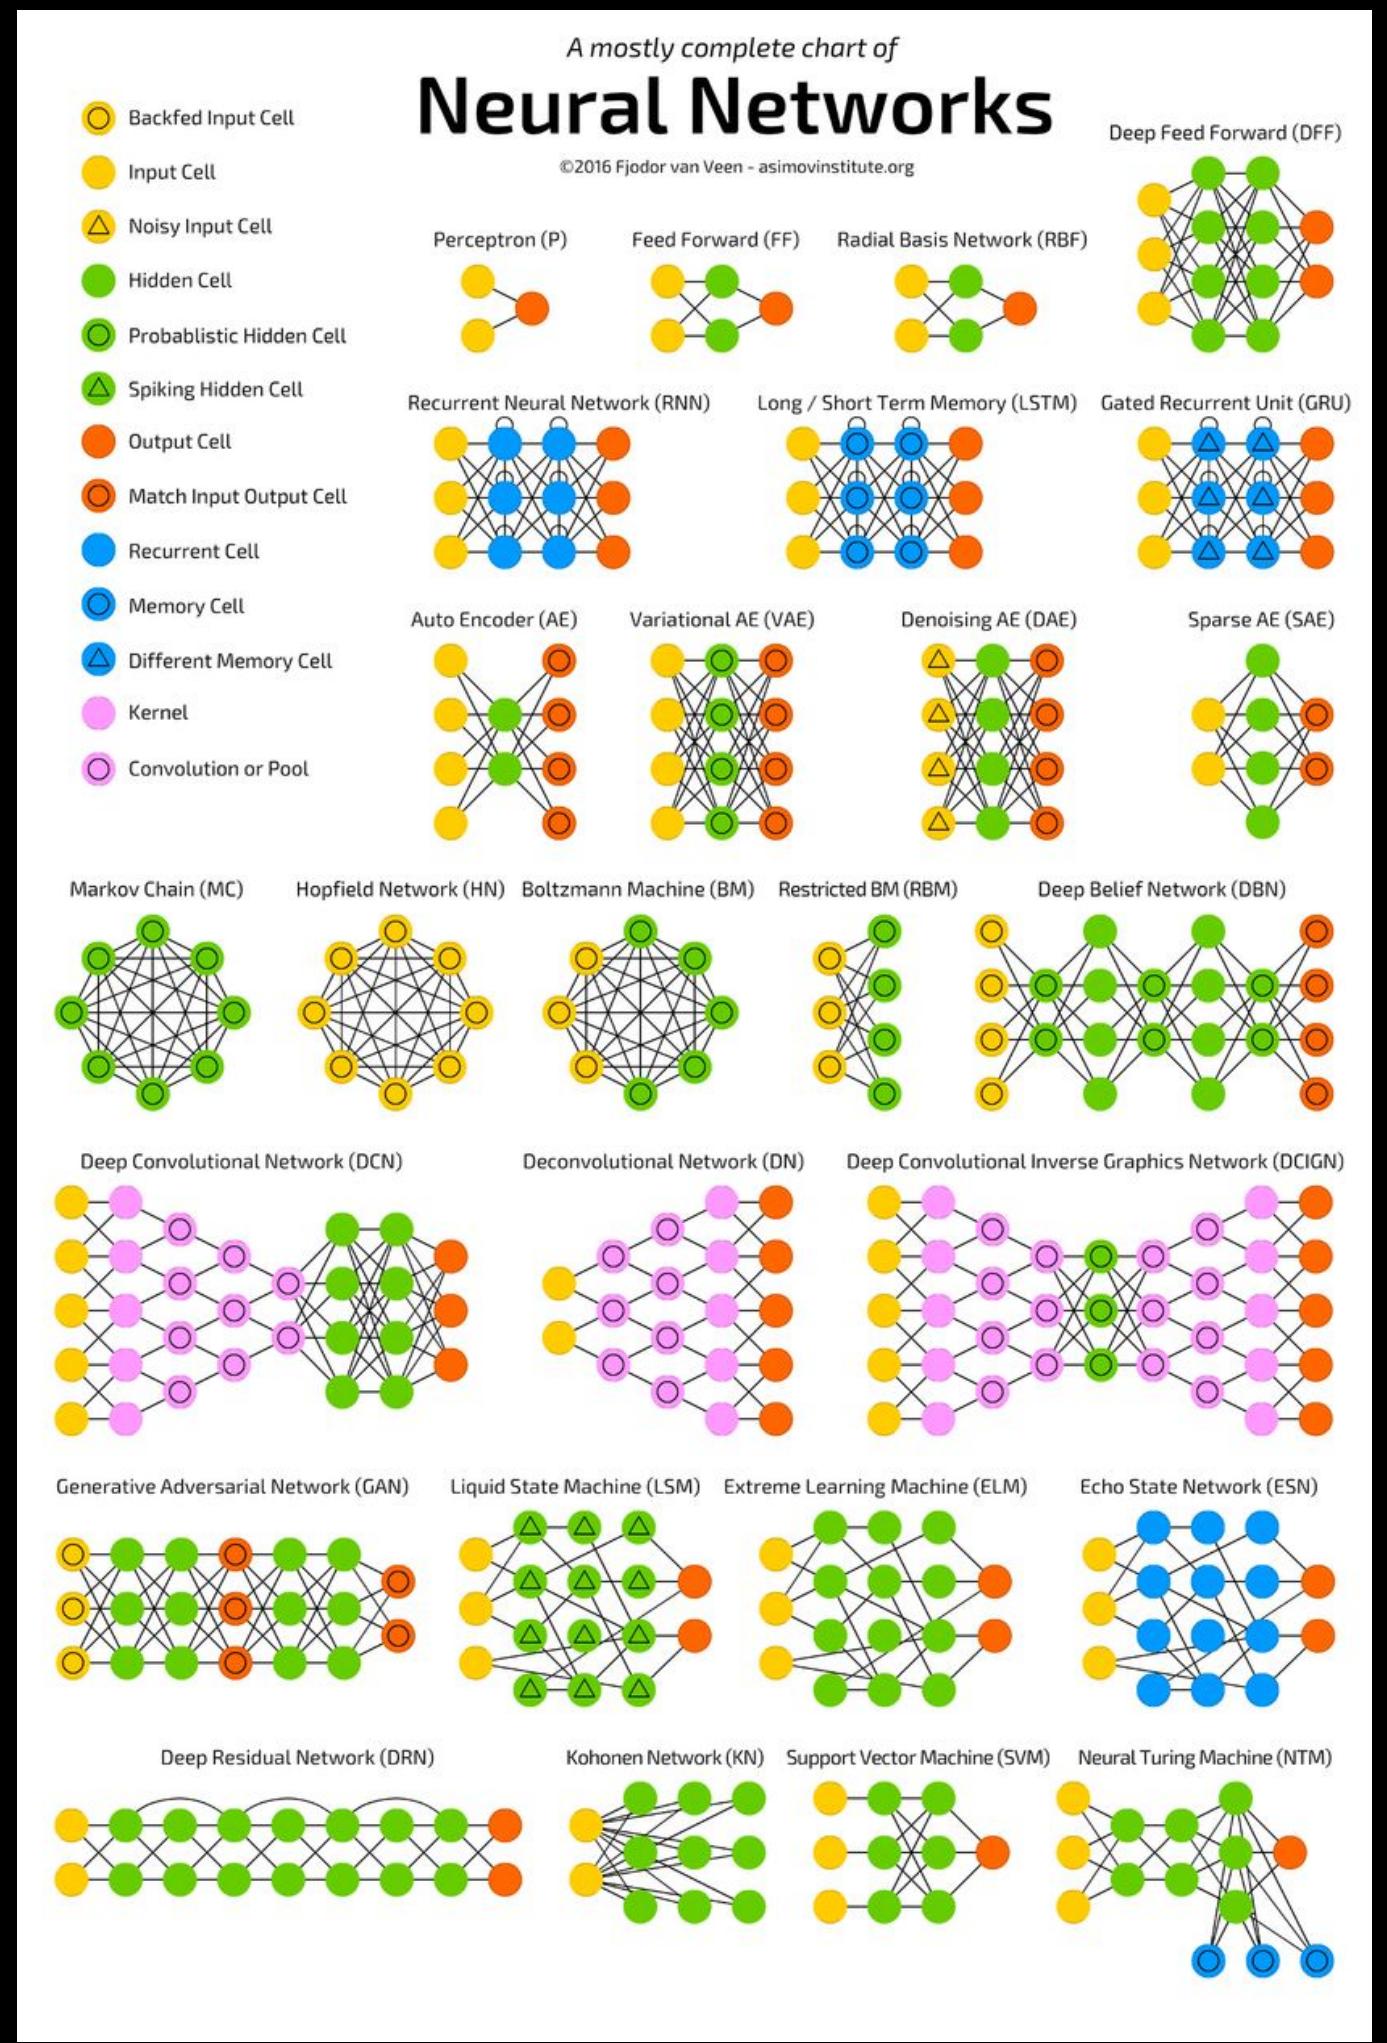
- Desactiva algunas neuronas (aleatoriamente) durante el entrenamiento
- << overfitting



blue = val
green = training

EARLY STOPPING

- Para el entrenamiento antes de llegar al overfitting



Optimización de hiperparametros

- Elegir el set optimo de hiperparametros para un algoritmo de aprendizaje (maximiza la performance del modelo)
- Son seteados antes del proceso de aprendizaje (#neurons, #cells, loss optimization algorithms, etc)

GRID SEARCH

- 1 — Identificar los hiperparametros del modelo a optimizar.
- 2 — Estimar el error obtenido para cada combinación en una grilla de hiperparametros.
- 3 — Selecciona la combinacion de hiperparametros con menor error

TRY THEM ALL



Optimización de hiperparámetros

```
# learning_rate choices
learning_rates = [ 0.1, 0.2, 0.3, 0.4, 0.5,
                    0.01, 0.02, 0.03, 0.04, 0.05 ]
# iterations choices
iterations = [ 100, 200, 300, 400, 500 ]

parameters = []
for i in learning_rates :
    for j in iterations :
        parameters.append( ( i, j ) )

print("Available combinations : ", parameters)

# Applying linear searching in list of available combination
# to achieve maximum accuracy on CV set

for k in range( len( parameters ) ) :
    # model = METHOD(..., learning_rate = parameters[k][0], iterations = parameters[k][1] )
    # ...
```

TRY THEM ALL



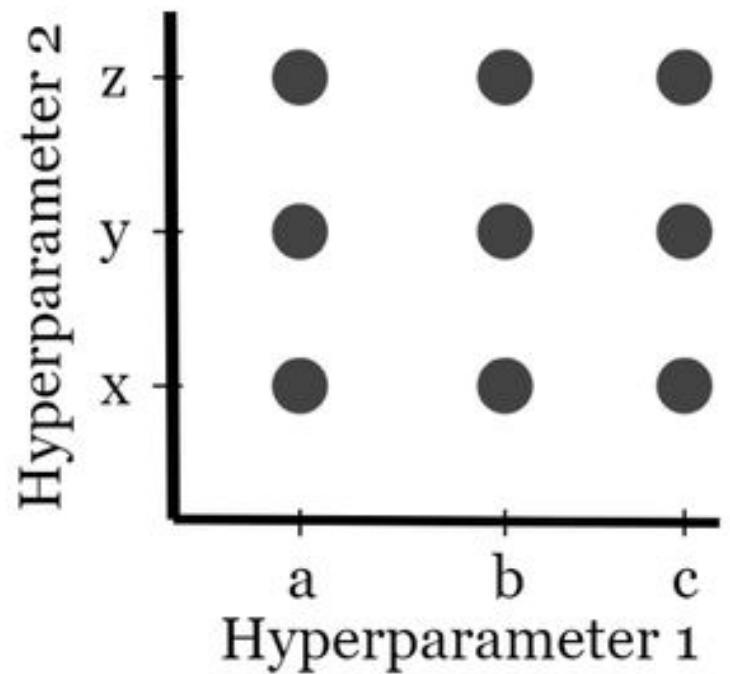
=GRID SEARCH

memegenerator.net

Optimización de hiperparámetros

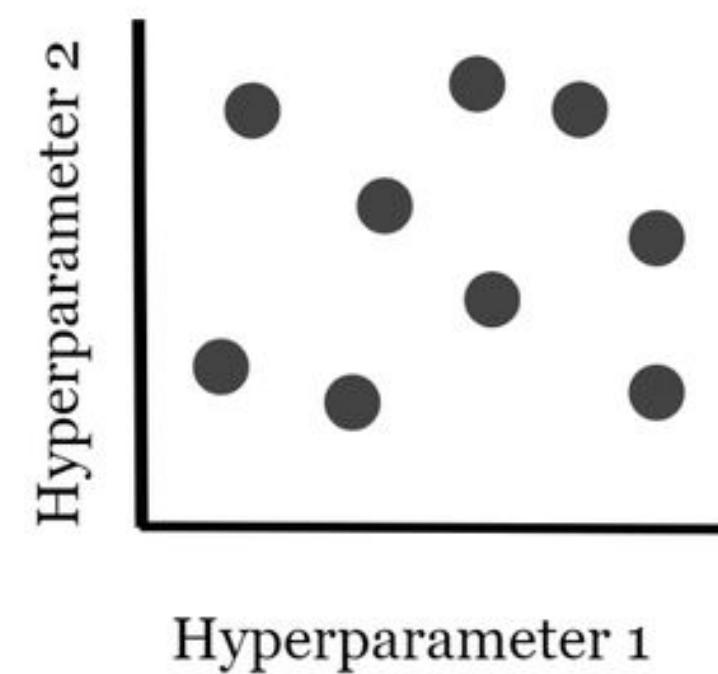
Grid Search

Pseudocode
Hyperparameter_One = [a, b, c]
Hyperparameter_Two = [x, y, z]



Random Search

Pseudocode
Hyperparameter_One = random.num(range)
Hyperparameter_Two = random.num(range)



Otros métodos

- Optimización Bayesiana
- Optimización Evolutiva

PROBLEMA-

cClasificar los ecos de un radar ionosférico segun dos clases (“good” or “bad”).
Hay 17 pares de números que representan 17 valores discretos correspondientes a la parte entera y compleja de la senial y que seran las entradas a la red neuronal o

VINCENT G. SIGILLITO, SIMON P. WING, LARRIE V. HUTTON, and KILE B. BAKER

CLASSIFICATION OF RADAR RETURNS FROM THE IONOSPHERE USING NEURAL NETWORKS

In ionospheric research, we must classify radar returns from the ionosphere as either suitable for further analysis or not. This time-consuming task has typically required human intervention. We tested several different feedforward neural networks to investigate the effects of network type (single-layer versus multilayer) and number of hidden nodes upon performance. As expected, the multilayer feedforward networks (MLFN's) outperformed the single-layer networks, achieving 100% accuracy on the training set and up to 98% accuracy on the testing set. Comparable figures for the single-layer networks were 94.5% and 92%, respectively. When measures of sensitivity, specificity, and proportion of variance accounted for by the model are considered, the superiority of the MLFN's over the single-layer networks is even more striking.

Paper source:

Sigillito, V. G., Wing, S. P., Hutton, L. V., \& Baker, K. B. (1989). Classification of radar returns from the ionosphere using neural networks. Johns Hopkins APL Technical Digest, 10, 262-266.

Data source:

Space Physics Group of the Johns Hopkins University Applied Physics Laboratory.

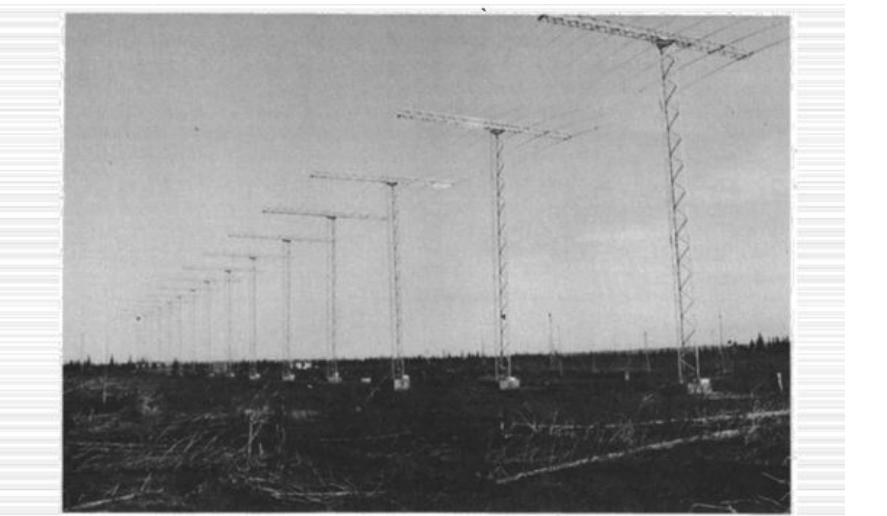
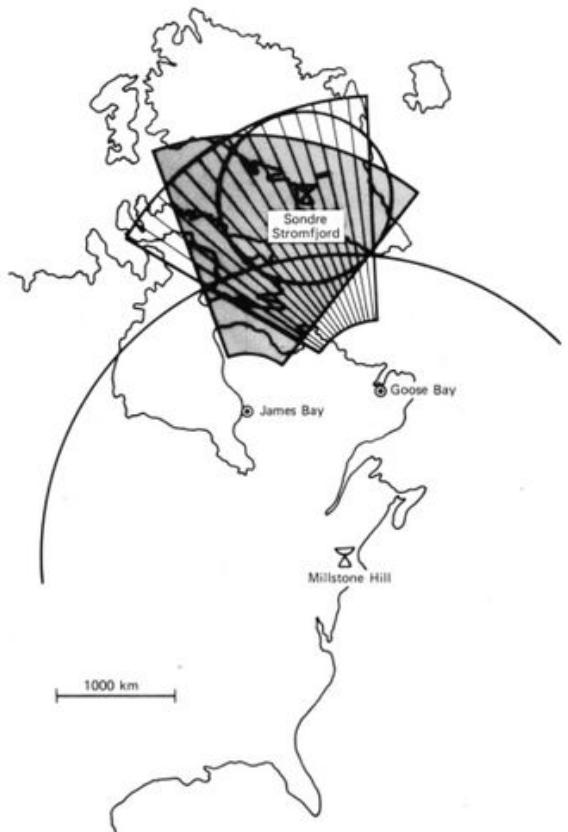
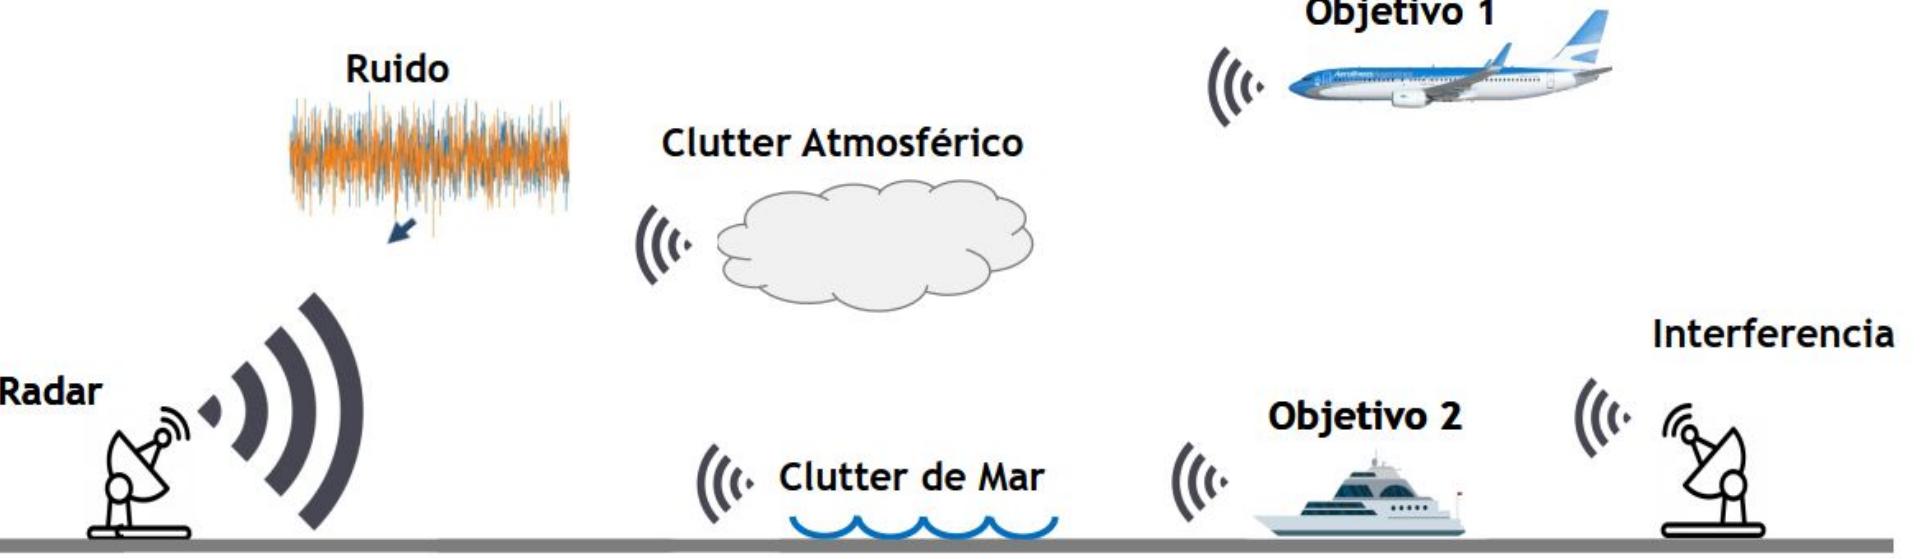
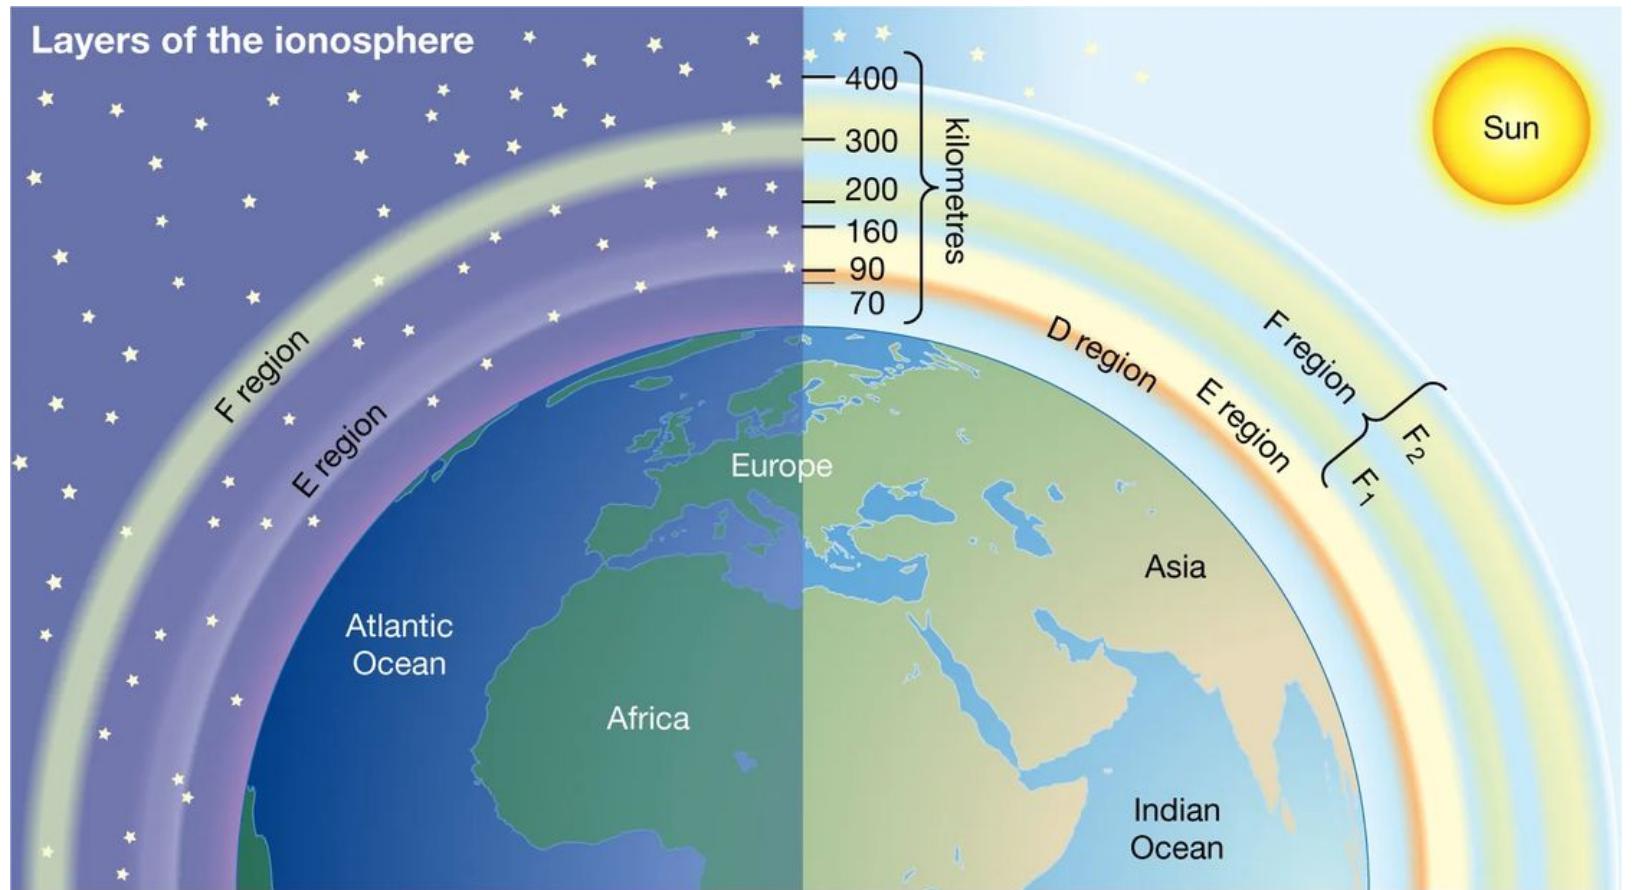


Fig. 7. View of antenna array from front, northwest corner. Shed for phasing matrix is located at the center of the array.

Dr. Z. Saavedra - FACET,
UNT



$$Eco(t) = Atte \cdot RCS \cdot u(t) \cdot m(t) \cdot e^{j\omega_ct \pm j\omega_D t + \varphi}$$

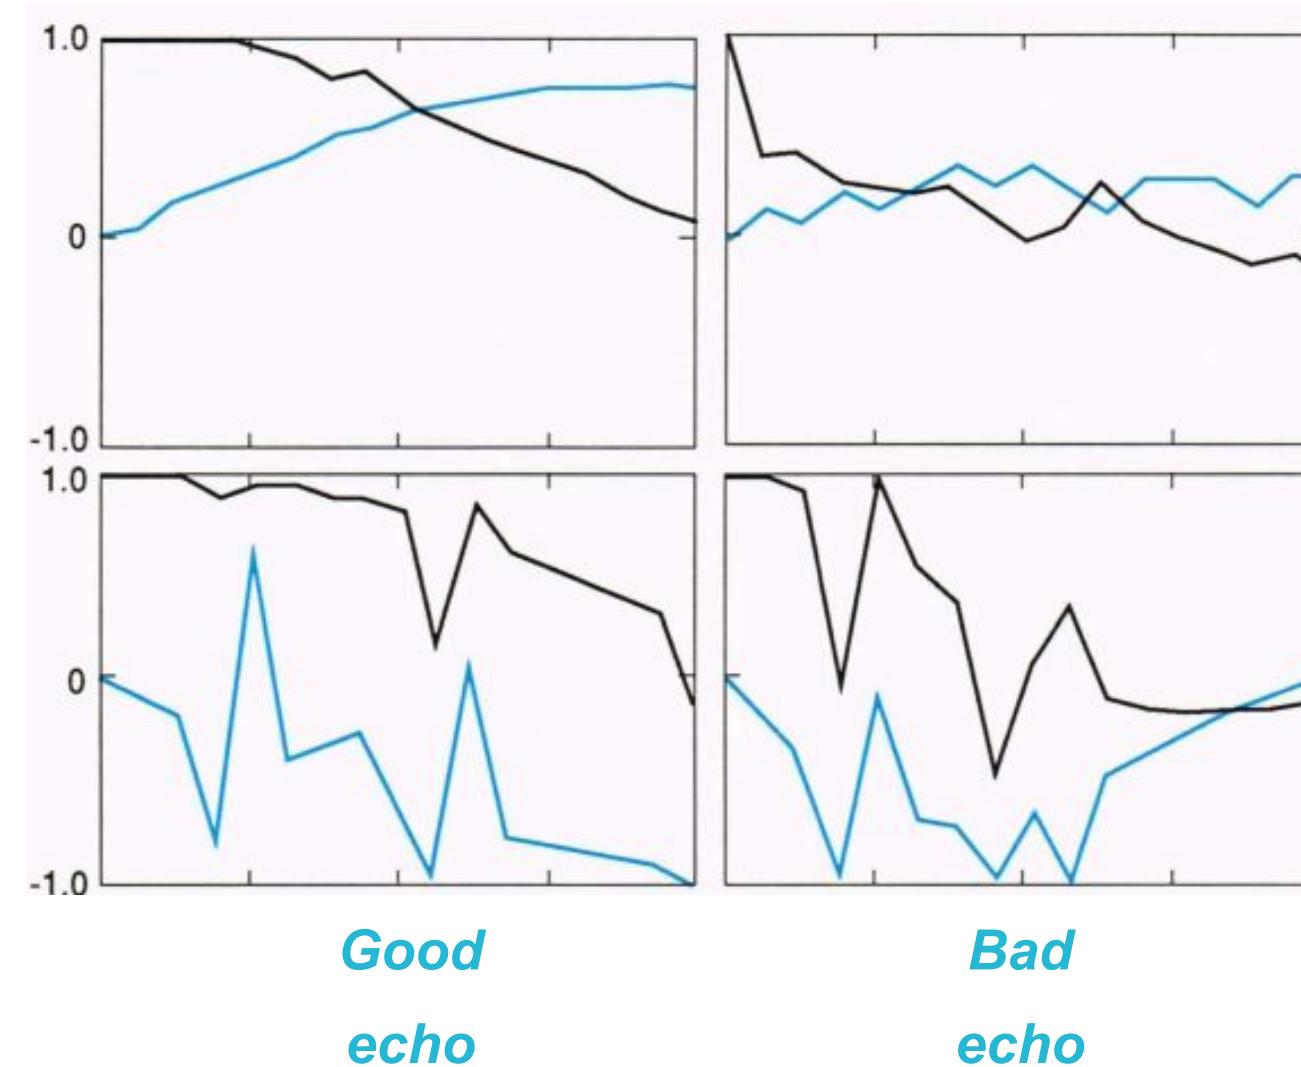


$$S_R(t) = Eco(t) + Clutter(t) + Ruido(t) + Interferencia(t)$$

*: "target" = ionosphere

IQ Modulation: it is a specific Phase-Modulation type, that provides us certain information in the real component of the signal and another possible coded information in the imaginary part. We can obtain both of them using demodulation techniques such as quadrature hybrids.

$$C(t) = A(t) + iB(t)$$



Auto-Correlation Function Output

